

Persistent Case State for Agentic Legal Analysis: A Structured Approach to Complex Queries

Sigvard Bratlie

April 14, 2026

Contents

1	Introduction	1
2	Theory	3
2.1	Large Language Models and Context Windows	3
2.2	Retrieval-Augmented Generation	3
2.3	Limitations of Naive RAG for Complex Reasoning	4
2.3.1	Semantic Similarity Does Not Guarantee Relevance	4
2.3.2	Cross-Document and Multi-Hop Reasoning	4
2.3.3	Unstructured Retrieval and Missing Information	4
2.4	Agentic Context Management	4
2.5	Related Work	5
2.5.1	Retrieval-Augmented Generation for Legal QA	5
2.5.2	Structured Extraction and Knowledge Organisation	5
2.5.3	Agentic Memory Management	6
2.5.4	Hierarchical and Abstractive Document Summarisation	6
2.5.5	Research Gap	7
3	Method	8
3.1	System Architecture	8
3.1.1	Overview	8
3.1.2	Agent Layer	8
3.1.3	Tool Layer	9
3.1.4	Project Structure	9
3.1.5	Data and Storage Layer	10
3.1.6	FactSheet Schema	10
3.1.7	Init and Update Pipeline	12
3.1.8	Clean Pipeline	13
3.1.9	LLM Configuration	14
3.1.10	Tech Stack Summary	14
3.2	Design Choices	15
3.2.1	Agent Layer and Orchestration	15
3.2.2	Database and Storage Architecture	15
3.2.3	FastAPI and Server-Sent Events	16
3.2.4	Concurrency, Reliability, and Tool Efficiency	16
3.3	Evaluation Design	18
3.3.1	Datasets	18
3.3.2	Evaluation Limitations and Caveats	21
3.3.3	Evaluation Environment	22
3.3.4	Metrics	23
3.3.5	Research Question Operationalization	24
3.3.6	Normality Testing	25
3.3.7	One-sided Hypothesis Tests	25
3.3.8	Levene's Test for Equality of Variance	26
3.4	Development Tooling and AI Assistance	27
3.4.1	Version Control and Repository Structure	27

3.4.2	Dependency Management	27
3.4.3	Report Authoring	27
3.4.4	Data Viewer	27
3.4.5	Use of AI Assistance	27
4	Results	29
4.1	Quality Metrics	29
4.1.1	DeepEval: Correctness, Relevancy, and Completeness	29
4.1.2	Query-Type Breakdown	34
4.1.3	BERTScore	35
4.1.4	Resource Metrics	37
4.2	Statistical Modeling	43
4.2.1	Mann-Whitney U Test	43
4.2.2	Levene's Test for Equality of Variance	44
4.2.3	Self-consistency	45
4.3	Summary by Research Question	45
5	Discussion	47
5.1	Research Question Findings	47
5.2	Value Beyond Performance Metrics	48
5.3	Uncertainties and Confounding Factors	49
5.4	Overall Assessment	50
6	Conclusion	51
6.1	Answers to Research Questions	51
6.2	Additional Finding: Variance Reduction	52
6.3	Contributions	52
6.4	Limitations and Future Work	52

Abstract

Large language models (LLMs) are increasingly applied to legal case analysis, where a single matter may span hundreds of pages across contracts, correspondence, expert reports, and court filings. The standard approach — retrieval-augmented generation (RAG) — addresses context size constraints by retrieving semantically similar document chunks, but fails on queries that require synthesizing facts distributed across multiple documents. This thesis proposes and evaluates *structured extraction* as a complementary approach: as case documents are ingested, key entities — parties, events, claims, damages, and deadlines — are extracted into a persistent, structured *FactSheet* that supplements chunk-based retrieval for case-level queries.

A LangGraph-based conversational agent implementing this architecture is evaluated against a RAG-only baseline across two legal datasets, multiple session stages, and three model configurations. Results show statistically significant improvements in correctness ($p = 8.5 \times 10^{-3}$) and relevancy ($p = 3.4 \times 10^{-4}$) without a proportional increase in token consumption or inference latency. A consistent and statistically robust reduction in score variance is also observed across all five evaluation metrics, suggesting that structured state reduces the brittleness introduced by retrieval failures. The performance gap is most pronounced for smaller, less capable models, indicating a practical cost-quality substitution: a cheaper model paired with the structured architecture can approximate the performance of a larger model on RAG alone. These findings demonstrate that schema-driven extraction is an effective and cost-efficient strategy for improving factual accuracy and reliability in multi-document legal settings.

Chapter 1

Introduction

The legal industry is increasingly adopting large language models (LLMs) to assist lawyers in navigating the complexity of modern litigation. Legal cases are inherently multi-document affairs: a single matter may involve contracts, correspondence, expert reports, court filings, and evidence spanning hundreds of pages. As a case develops, new attachments arrive, earlier positions are revised, and the relevance of existing documents shifts. The central challenge for an LLM-based legal assistant is not merely reading these documents — it is maintaining a coherent, accurate understanding of the case as a whole, across all its documents and all its turns of development.

A foundational constraint of current LLM architectures makes this challenge particularly acute: inference endpoints are stateless. Each call to the model is an independent, context-free computation — the model retains no knowledge of prior interactions, prior documents, or prior reasoning once a response is returned. Whatever the agent “knows” about a case must be reconstructed and injected anew at every query, from conversation history, retrieved chunks, or some other persistent external representation. For a legal assistant that must maintain a coherent picture of a case across days, sessions, and an evolving document set, statelessness means that memory is not a default capability — it is a design problem. The system must explicitly decide what to remember, where to store it, and how to surface it efficiently at query time.

The standard approach to this problem is retrieval-augmented generation (RAG): rather than loading all documents into context, the system retrieves the chunks most semantically similar to the user’s query. This is a practical solution to the context size problem, but it introduces a different and more fundamental limitation. Legal questions frequently require synthesizing facts that are distributed across multiple documents — a party’s identity established in one contract, a disputed event described in correspondence, a damages figure appearing in an expert report. Chunk-based retrieval cannot produce this synthesis: it operates on isolated local fragments, and high semantic similarity to a query does not guarantee that the retrieved chunks contain the relational context needed to answer it correctly.

An alternative — injecting all documents directly into a large context window — does not solve the problem either. Even as model context windows have grown to one million tokens and beyond, empirical evidence shows that effective utilization degrades well before the nominal limit, and model performance deteriorates for information buried deep in a long context. Beyond capacity, loading an entire case file for every query is impractical and costly. The goal should instead be to provide the minimum amount of information that maximizes relevance to the query at hand.

This thesis proposes and evaluates a third approach: *structured extraction*. As case documents are ingested, key facts — parties, claims, events, damages, deadlines, etc. — are extracted and organized into a structured *FactSheet*: a persistent, machine-readable summary of the case. The FactSheet supplements chunk-based RAG for queries that require a case-level overview, providing relational context that retrieval alone cannot assemble. It is initialized at project creation and updated incrementally as new documents arrive, so the agent always has access to a coherent picture of the case without re-reading the full document history.

This thesis addresses the following research questions:

Main RQ: *Does schema-driven cross-document extraction into a structured case representation improve the factual accuracy of an agentic RAG system in multi-document legal settings, without incurring*

proportional increases in query-time cost?

RQ1: *To what extent does structured extraction mitigate the cross-document retrieval limitations of naive RAG in legal case analysis?*

RQ2: *Does the structured case representation maintain or improve agent performance as document volume grows across case sessions?*

RQ3: *Can a structured extraction architecture compensate for reduced model capacity — enabling a smaller, less expensive model to achieve performance comparable to a larger baseline?*

The system is implemented as a LangGraph-based conversational agent. The main RQ is evaluated through a controlled comparison of a structured-state agent against a RAG-only baseline across factual quality and resource consumption metrics. RQ1 is addressed by examining performance on queries that require cross-document synthesis. RQ2 is examined through session-level trajectories as document volume grows. RQ3 is addressed by comparing performance gaps across models of varying capability.

The contributions of this thesis are: (1) the design and implementation of initialization and incremental update pipelines for structured legal case state management, (2) an evaluation framework comparing structured state agents against RAG-only and baseline approaches across multiple models and datasets, and (3) an empirical evaluation of how structured extraction affects performance across datasets, model configurations, and document volumes.

The remaining chapters are organized as follows: Chapter 2 presents the theoretical background on LLM context limitations, RAG and its limitations, and structured extraction approaches. Chapter 3 describes the methodology, system architecture, and evaluation design. Chapter 4 presents experimental results. Chapter 5 discusses findings and implications, and Chapter 6 concludes with a summary and directions for future work.

Chapter 2

Theory

This chapter builds the theoretical case for the approach taken in this thesis. The argument unfolds in two movements: first, why large language models alone are insufficient for multi-document legal case analysis; and second, why retrieval-augmented generation — the standard remedy — carries fundamental limitations that make it inadequate for cross-document reasoning. These limitations motivate the structured extraction approach evaluated in this thesis. Section 2.5 then surveys prior work and alternative architectures that contextualise the design choices.

2.1 Large Language Models and Context Windows

Modern LLMs are built on the transformer architecture [1], which processes input as a sequence of tokens bounded by a fixed context window. Recent models advertise increasingly large context windows — some exceeding one million tokens — which may appear to solve the problem of large document corpora. This appearance is misleading for at least three reasons.

First, the nominal context size overstates effective capacity. Hsieh et al. introduce RULER, a benchmark measuring the *real* usable context of long-context LLMs, and demonstrate that performance degrades substantially before the advertised limit is reached [2]. Second, LLMs do not attend uniformly across the context: performance is highest when relevant information is at the beginning or end, and degrades for information buried in the middle — the *lost in the middle* phenomenon [3]. Third, and most practically, a legal case involving dozens of attachments spanning hundreds of pages cannot be loaded in full for every query without prohibitive cost and latency, regardless of model capability.

The implication is a design principle: rather than maximizing the amount of text in context, the goal should be to include the *minimum* amount of information while ensuring that information is as relevant as possible to the query at hand. This principle rules out naive full-document injection and motivates retrieval as the primary mechanism for selecting what enters context. The question then becomes: what are the limits of retrieval?

2.2 Retrieval-Augmented Generation

Retrieval-Augmented Generation (RAG) operationalizes the minimum-relevant-information principle by retrieving only the most semantically similar document chunks at query time, rather than loading entire documents [4]. Documents are chunked and embedded offline into a vector store; at inference time, the top- k most similar chunks are retrieved and inserted into the LLM context alongside the query.

For straightforward factual lookup — “what is the defendant’s registered address?” — this works well. The answer lives in a single location in the document corpus, it is semantically similar to the query, and a chunk-based retriever will find it. The problems emerge when queries require understanding that spans multiple documents, or when the relevant information is not directly expressible as a similarity to the query string itself.

2.3 Limitations of Naive RAG for Complex Reasoning

2.3.1 Semantic Similarity Does Not Guarantee Relevance

The retrieval step selects chunks based on embedding similarity to the query. High similarity does not guarantee true relevance: documents may share keywords or surface features without addressing the underlying intent or logic of the query [5]. In legal contexts, where questions about causation, liability, or damages require understanding relational chains across documents, surface-level matching is an insufficient retrieval criterion.

Compounding this, injecting all retrieved chunks directly into context can produce fragmented or contradictory input that confuses the LLM and degrades answer quality [5]. The problem is not just what is retrieved but how it combines in context.

2.3.2 Cross-Document and Multi-Hop Reasoning

The more fundamental limitation is that the answer to a legal query often does not reside in a few isolated chunks, but is distributed across multiple documents. Existing naive RAG systems are inadequate for queries that require retrieval and reasoning over multiple pieces of supporting evidence across documents [6]. In multi-document settings, relevant information is spread across sources and requires coherent synthesis: naively concatenating retrieved passages leads to fragmented or contradictory responses, a problem that is especially acute in legal and domain-specific QA where multi-hop reasoning is critical [7].

This limitation is structural, not a matter of tuning. A chunk embedding represents a local fragment of a single document. No similarity threshold or retrieval hyperparameter can make a chunk-based retriever produce a coherent synthesis of facts distributed across an entire case file. The apparent alternative — injecting the full document corpus into context — was ruled out in Section 2.1: effective utilisation degrades well before the nominal context limit is reached, cost and latency scale prohibitively with document volume, and the design principle established there requires minimising rather than maximising the information loaded at query time.

2.3.3 Unstructured Retrieval and Missing Information

Standard RAG processes unstructured text passages without leveraging structured domain knowledge, and struggles with complex queries requiring multi-hop reasoning or domain-specific knowledge organization [8]. Unstructured retrieval also risks returning passages that contain non-atomic information — facts entangled with irrelevant surrounding text — that can mislead the LLM rather than inform it [8]. Chunk size amplifies this: small chunks provide insufficient context for understanding a passage, while large chunks dilute the semantic signal and retrieve off-topic content.

These three limitations — relevance mismatch, cross-document reasoning failure, and unstructured knowledge gaps — together define the core challenge this thesis addresses. A legal agent that relies solely on chunk-based RAG will, by design, struggle with case-level questions that require a coherent picture of the whole. This motivates approaches that move beyond isolated chunk retrieval toward structured, relational representations of case knowledge.

2.4 Agentic Context Management

The final theoretical component concerns how structured knowledge is maintained across sessions in a stateless architecture. As noted in the introduction, LLM inference endpoints carry no state between calls: context must be reconstructed at every query. Stateful agent frameworks such as LangGraph address this through checkpointing — persisting structured agent state to a database and reinjecting it at query time [9]. For the conversational history specifically, rolling summarization — periodically condensing prior turns into a compact summary — is a practical strategy for preventing context overflow [10]. It is worth noting that rolling summarization applies to the chat inference pipeline only; the document ingestion pipeline operates as a stateless batch process and does not maintain rolling conversational context.

In the system developed in this thesis, the factsheet is the primary persistent state artifact. It is initialized during document ingestion and updated incrementally as new attachments are processed

and new facts confirmed, providing the agent with a structured, token-efficient overview of the case at every point in the conversation. The theoretical argument, then, is that this persistent structure addresses what RAG cannot: a coherent, relational picture of the case that is available without vector-based retrieval: the factsheet is loaded directly from the structured database, not assembled from semantically matched chunks at query time.

The following chapter describes how these theoretical motivations are translated into system design and how their combined effect is evaluated empirically.

2.5 Related Work

This section surveys prior work on the core problems this thesis addresses: overcoming RAG limitations in large multi-document corpora, structured information extraction for complex QA, and agentic LLM systems for professional knowledge work. The legal domain is the application context, but the architectural challenges — handling large, heterogeneous document collections with distributed factual content — are not unique to law. Related problems appear in financial due diligence, medical records analysis, and other high-stakes professional domains where information is voluminous, structured, and spread across many sources.

It is worth noting that most prior work on LLMs in legal settings concerns legal *reasoning* — statutory interpretation, case law analysis, and normative prediction [11, 12]. This thesis is not primarily concerned with legal reasoning ability. The focus is on the prior, more practical problem: how to make the factual content of a large case file reliably accessible to an LLM agent, regardless of which model underlies it. This distinction matters when interpreting the related work: findings about legal language models and legal reasoning benchmarks address a different problem, and this thesis fills a gap that those works do not cover.

2.5.1 Retrieval-Augmented Generation for Legal QA

RAG has emerged as the standard architecture for domain-specific document QA, and its application to legal settings has been explored in several directions.

Louis et al. [13] demonstrate that a RAG-augmented LLM pipeline can generate long-form answers to statutory law questions with promising automatic metric scores, while incorporating extractive rationales for improved interpretability and citation grounding. Their system retrieves relevant statutory passages (but not case law) and incorporates them as context at inference time — the same basic mechanism used by the baseline agent in this thesis. Their work highlights both the strengths of RAG (traceable sourcing, factual grounding) and its limitations when queries require synthesising evidence from across a large corpus or when hallucinations persist despite relevant retrieval.

The fundamental tension identified in RAG-based legal systems is that legal questions rarely decompose into a single, retrievable passage. Multi-hop reasoning over distributed evidence is common: establishing liability may require linking a party’s obligation (in a contract), its breach (in correspondence), and the resulting damage (in a financial report), none of which are co-located in the document corpus. MultiHopRAG [6] and HopRAG [14] address this limitation in general QA settings by explicitly chaining retrieval steps, constructing logical paths across sources. This thesis takes a complementary approach: rather than chaining retrievals reactively at query time, it extracts and pre-structures the key relational facts during document ingestion, making them available as a persistent case state without requiring multi-hop retrieval.

2.5.2 Structured Extraction and Knowledge Organisation

Extracting structured information from unstructured legal documents — parties, events, claims, and temporal relationships — has been studied as an information extraction problem long before the LLM era. Recent work has revisited this problem with LLMs: Jiang et al. [8] propose Retrieval and Structuring Augmented Generation (RSAG), which combines retrieval with LLM-driven structuring of retrieved content before generation. Their results show that structuring retrieved passages into typed entities before injecting them into the generation context consistently outperforms naive RAG on complex reasoning tasks. This provides independent empirical support for the core design hypothesis of this thesis: that structured, relational case state is more useful to an LLM than raw retrieved chunks.

GraphRAG [15] extends this idea to the full document corpus, building a knowledge graph over entities and their relationships to enable community-level summarization and structured traversal. HopRAG [14] adds explicit multi-hop reasoning, constructing logical chains across documents rather than relying on chunk similarity alone. The factsheet implemented in this thesis can be understood as a lightweight, domain-specific instance of the same principle: rather than constructing a general-purpose property graph, it organises only the entities most relevant to legal case management into a typed, relational structure. The reduced generality is an acknowledged trade-off — the system covers only what a legal practitioner is likely to ask rather than supporting arbitrary queries over the full document space. The practical advantages of this choice, in terms of query-time cost, schema precision, and interpretability, are examined in the evaluation and discussion.

2.5.3 Agentic Memory Management

A structural constraint shared by all retrieval-based approaches is that the LLM context window is finite. MemGPT [16] addresses this directly by treating the context window as a fixed-capacity “main context” — analogous to RAM — and introducing a tiered memory architecture in which an explicit memory manager decides what information to page in from external storage and what to evict. Rather than relying on a single retrieval query, the system supports fine-grained memory operations: appending new facts to context, archiving stale content to external storage, and selectively restoring it on demand. The result is an agent that can maintain coherent state across long interactions without ever exceeding the physical context limit.

Several of these mechanisms inform the design of the system in this thesis. The factsheet functions as a domain-specific persistent memory: it is initialised during document ingestion and injected directly into context at query time, bypassing the need for vector retrieval of case-level facts. In the chat inference pipeline, rolling summarisation of conversation history serves an analogous eviction function, preventing accumulated dialogue from consuming the context budget; the document ingestion pipeline, being a stateless batch process, does not require this mechanism.

A complementary diagnostic perspective is offered by Yuan et al. [17], who decompose agent memory failures into two distinct bottlenecks: *retrieval failures*, in which the correct information is not surfaced from storage, and *utilisation failures*, in which the correct information is retrieved but not correctly used by the model. Their empirical analysis shows that improving retrieval quality consistently yields larger performance gains than increasing the sophistication of the write-time encoding process. This finding supports the design priority in this thesis: the factsheet eliminates retrieval entirely for case-level facts by placing them directly in context, targeting the retrieval bottleneck rather than investing in more elaborate embedding strategies.

2.5.4 Hierarchical and Abstractive Document Summarisation

RAPTOR [18] takes a different approach to the multi-document comprehension problem. Rather than retrieving isolated chunks, it recursively clusters document segments and generates abstractive summaries at each level of the cluster hierarchy, constructing a tree whose leaves are raw text chunks and whose internal nodes are progressively higher-level summaries. At query time, retrieval can target any level of the tree, enabling both fine-grained lookup and broad, thematic search that spans many documents.

RAPTOR directly addresses the cross-document reasoning failure identified in Section 2.3: because higher-level nodes aggregate content from many source segments, a single retrieval step can surface information that would otherwise require multiple hops across the corpus. In practice, however, the approach carries significant costs. Constructing the hierarchy requires summarising every cluster at every level, which is computationally expensive and heavily dependent on the quality of the summarisation model. Weak summarisation propagates errors upward through the tree, and the cost scales with both corpus size and model capability. Furthermore, the summaries produced by RAPTOR are consumed solely as retrieval targets: they improve the chances that a relevant passage is found, but they are not themselves injected as structured context.

The factsheet approach taken in this thesis can be understood as a more targeted variant of the same intuition. Rather than abstracting the entire corpus into a general-purpose hierarchy, it extracts only the entities most relevant to legal case management — parties, claims, damages, key events, deadlines, etc — and organises them into a typed, relational structure. This structure serves a dual

purpose: it both enables structured retrieval of specific case facts and is injected directly as context, giving the model a coherent case overview without any retrieval step. The trade-off is scope: the factsheet covers only what a legal practitioner is likely to ask, rather than supporting arbitrary queries over the full document space.

2.5.5 Research Gap

Taken together, the related work establishes several important results: (1) general RAG is insufficient for multi-document legal reasoning due to relevance mismatch and cross-document synthesis failures; (2) structured extraction consistently improves performance on complex QA tasks; (3) graph-based and multi-hop retrieval partially address cross-document reasoning by making entity relationships explicit; (4) tiered memory management, as demonstrated by MemGPT, enables coherent long-horizon agent behaviour by selectively controlling what enters context; and (5) hierarchical abstractive summarisation, as in RAPTOR, improves broad-coverage retrieval but at significant computational cost and without producing a directly injectable, domain-typed representation.

What none of these lines of work address is the combination: a domain-specific, persistently maintained structured case state that is injected directly into context — bypassing retrieval for the information that matters most — within a conversational agent evaluated on authentic legal case files across multi-session workflows. Prior systems are either domain-agnostic (MemGPT, RAPTOR, GraphRAG), evaluated on synthetic or curated benchmarks rather than real case management tasks, or focused on legal reasoning rather than the practical problem of making large, heterogeneous case files reliably accessible. This thesis addresses that gap by implementing and empirically evaluating this architecture against a RAG baseline using real legal case documents and expert-constructed ground truth.

Chapter 3

Method

3.1 System Architecture

The system is organized as a `uv` workspace containing one core module and three auxiliary modules. The **agent** package is the primary contribution of this thesis: it implements the full backend, including the LangGraph agent, document processing pipelines, vector store integration, and database management. The remaining packages are supporting tools: **ui** provides a Streamlit interface for non-technical users to interact with the agent; **evals** contains the evaluation harness used to benchmark agent configurations; and **data-viewer** is a lightweight helper for inspecting and editing evaluation datasets. The architecture description that follows concerns the **agent** module exclusively.

3.1.1 Overview

The backend is implemented as a **FastAPI** [19] application (`agent/main.py`) that exposes a RESTful API with Server-Sent Events (SSE) for real-time response streaming. The frontend is a **Streamlit** application that communicates with the backend over HTTP, consuming streamed responses. Authentication is handled via **Supabase Auth** using JWT tokens, which are validated on each request through a `HTTPBearer` dependency in FastAPI.

3.1.2 Agent Layer

The core agent is implemented as a LangGraph **StateGraph** (see Figure 3.1). The graph has three nodes:

- **init** — Injects the system prompt as the first **SystemMessage** on a new conversation. When resuming an existing session, the system prompt is already present in the persisted checkpoint state and this node performs no action.
- **call_llm** — Assembles the full input payload (system prompt, factsheet context, conversation history, and any per-request attachment content), sanitizes it for provider compatibility, then streams the LLM response.
- **call_tool** — Executes all tool calls emitted by the LLM in the previous turn, with per-call caching to avoid redundant invocations within the same session.

Control flows as: **START** → **init** → **call_llm**. A conditional edge then routes to **call_tool** if the LLM response contains tool calls (looping back to **call_llm** afterwards), or to **END** if not. To manage long conversations within the model’s context window, a rolling summarization strategy is applied every **sum_rate** messages (configurable, default 20) using a dedicated **Summarizer** module.

Conversation state is persisted using LangGraph’s **AsyncPostgresSaver** checkpointing, backed by **Supabase PostgreSQL**, enabling session resumption and multi-turn dialogue across requests.




Figure 3.1: Agent graph: `init` → `call_llm` with a conditional loop through `call_tool`.

3.1.3 Tool Layer

The agent has access to a set of callable tools. The full tool set used in the primary agent configuration includes:

- **query_project_attachments** — Performs semantic retrieval (RAG) against project-specific document chunks indexed in the vector store.
- **show_elements** — Returns structured factsheet elements (events, parties, claims, damages, deadlines) for a project, optionally filtered by date range and significance level.
- **list_attachments** — Lists a project’s attachments and emails (with metadata such as title, date, and category), optionally filtered by date range and significance. Returns file IDs that the model can pass to **read_attachments**.
- **read_attachments** — Fetches and returns the full parsed content of one or more documents by file ID. The parsed body is cached in the PostgreSQL database (populated during the document analysis pipeline); if not yet cached, the raw file is read from GCS and parsed on demand.

There are also secondary tools that are not directly relevant for this thesis, but are included in both agents for completeness and future extensibility:

- **tavily_search** — Web search via the Tavily API.
- **query_laws** — Semantic retrieval against a corpus of Norwegian legislation, with optional filtering by law title.
- **read_specific_law** — Retrieves exact paragraph text from legislation by title and paragraph number via a parameterised BigQuery query.

An ablated tool configuration is defined for evaluation purposes. It includes the secondary tools and project document RAG (**query_project_attachments**), but excludes the structured factsheet tools and direct file-reading access.

3.1.4 Project Structure

The repository is organized as a `uv` workspace. The simplified directory tree below shows the primary modules and their responsibilities.

```

master-thesis/
|-- agent/                                # Core backend (primary contribution)
|   '-- src/
|       |-- agent/
|           |-- agent.py                  # Conversational LangGraph StateGraph
|           |-- pipelines.py              # Init/update document pipelines
|           |-- agent_modules.py          # Summarizer, ToolManager
|           |-- context_manager.py        # ContextManager (extraction, cleanup)
|           |-- tools.py                  # Tool definitions
|           '-- utils.py                  # Prompts, LLM factory
|       |-- database/
|           |-- database_modules.py        # SupabaseManager (CRUD)
|           '-- vectorstore_modules.py     # BQVectorStore, ChromaVectorStore
|       |-- documents/
|           '-- document_modules.py        # DocumentProcessor, EmailHandler
|       '-- models/
|           |-- project_models.py          # FactSheet, Party, Event, Claim, ...
|           '-- agent_models.py           # AgentState (LangGraph TypedDict)
|-- ui/                                    # Streamlit frontend
|-- evals/                                # Evaluation framework
|-- data-viewer/                          # Dataset inspection tool
'-- thesis/                               # This document

```

Figure 3.2: Simplified repository structure. The `agent` package is the primary contribution; the remaining packages are supporting tools.

3.1.5 Data and Storage Layer

Document persistence is handled through two complementary mechanisms. Raw files (PDF, DOCX, PPTX, EML) are stored in **Google Cloud Storage (GCS)**, where they can be retrieved on demand by the `read_attachments` tool when no parsed body cache is available. Parsed and chunked document content is embedded and stored in a persistent **vector store** for semantic retrieval. Structured project data — including the **FactSheet** and session metadata — is stored in a **PostgreSQL** database via the **SupabaseManager**, which provides create, read, update, and delete (CRUD) operations for all relational data.

3.1.6 FactSheet Schema

The *FactSheet* is the central persistent artifact of the system: a structured, machine-readable representation of the legal case that the agent maintains across the full lifecycle of a matter. It is stored relationally in PostgreSQL and loaded into context on each agent call, supplementing vector retrieval for case-level queries. The schema is defined using Pydantic **BaseModel** classes, which enforce field types and validation at runtime and serve as the structured output target for all LLM extraction calls.

The top-level **FactSheet** object aggregates five typed entity collections alongside a free-text background and case metadata. Table 3.1 summarises the complete schema, organised by entity type.

Table 3.1: FactSheet entity schema. All entity types share a **significance** field (*high* / *medium* / *low*) and a **file_id** / **email_id** back-reference to the source document.

Entity	Field	Description
<i>FactSheet</i>	title	Short case title (max 10 words)
	background	Narrative factual summary of the case
	start_date	Date the matter began
<i>Party</i>	legal_name	Full legal name of the party
	role	Functional role: <i>plaintiff</i> , <i>defendant</i> , <i>witness</i> , <i>expert</i> , <i>contractor</i> , etc.
	entity_type	<i>individual</i> , <i>company</i> , or <i>government</i>
	role_description	Free-text elaboration of the party's involvement
	party_reps	List of named legal representatives (name, role, contact)
<i>Event</i>	event_name	Short descriptive name
	event_start_date	Start date (partial dates coerced to first of month)
	event_end_date	End date if applicable
	description	Narrative description of what occurred
	category	e.g. <i>court_filing</i> , <i>contract_signing</i> , <i>communication</i>
	disputed	Boolean flag: is this event contested?
<i>Claim</i>	category	<i>principal claim</i> , <i>counterclaim</i> , <i>objection</i> , <i>ancillary claim</i> , etc.
	legal_basis	Statutory or contractual basis (e.g. <i>avtaleloven §36</i>)
	factual_basis	Key facts that support the claim
	relief_sought	What is being claimed (damages, injunction, etc.)
	strength_assessment	LLM assessment: <i>strong</i> , <i>moderate</i> , or <i>weak</i>
<i>Damage</i>	party_role	Which party asserts this claim
	category	<i>direct losses</i> , <i>interest</i> , <i>consequential</i> , or <i>punitive</i>
	amount	Monetary amount if stated; <i>None</i> if unspecified
	currency	e.g. <i>NOK</i> , <i>USD</i>
	basis	Factual basis for the damage claim
	supporting_evidence	List of file_ids that evidence this damage
<i>Deadline</i>	party_role	Which party asserts this damage
	deadline_date	The deadline date
	description	What the deadline concerns
	party_role	Which party is subject to this deadline

Every entity carries a **significance** field ("high", "medium", "low"), assigned by the extracting LLM based on the entity's centrality to the case. This field is exposed as an optional filter parameter in the factsheet tool calls: when invoking **show_elements**, the agent can request only entities at or above a given significance level, retrieving a focused subset rather than the full collection. During standard context injection — when the factsheet is loaded directly as part of the system prompt at the start of each inference call — all entities are included regardless of significance. Each entity also carries a **file_id** or **email_id** back-reference to the source document, enabling the agent to trace any extracted fact back to its evidentiary origin.

The schema is deliberately domain-specific: rather than a general-purpose knowledge graph, it captures precisely the entity types that legal practitioners reason over in a case management context — who the parties are, what happened and when, what is claimed and on what legal basis, what damages are asserted, and what procedural deadlines apply. This specificity is a design trade-off (discussed further in Section 3.2): the schema matches the query patterns of legal practitioners but requires adaptation for other legal domains.

3.1.7 Init and Update Pipeline

Project initialization (`POST /init-project`) and updates (`POST /update-project`) are handled by the `ProjectPipeline` class, which compiles a `LangGraph StateGraph` for each operation (see Figure 3.3). The graph exploits `LangGraph`'s fan-out semantics to run independent branches in parallel, minimising total wall-clock time for the ingestion of large document sets.

Initialization pipeline. The following describes each node and its role in the `init` graph, in execution order:

- `collapse_emails` and `initialize_input` are launched concurrently from `START`. `collapse_emails` deduplicates EML threads by collapsing reply chains into canonical thread objects, preventing the same email content from being processed multiple times. `initialize_input` runs a lightweight LLM call over the user-provided case description to produce an initial structured seed for the factsheet: a short title, a factual background summary, and a list of named parties. This seed is used immediately as the factsheet starting state before any documents are processed.
- `extract_emails` follows `collapse_emails` and decodes each collapsed email thread, extracting inline attachments and populating the pipeline's document list with any files referenced in the email bodies.
- `parsing` and `storage` both fan out from `extract_emails` and run concurrently. `parsing` processes each document into plain text: PDFs are parsed with `PyPDF2` (with OCR via `ocrmypdf` as fallback for scanned documents), DOCX files with `python-docx`, PPTX with `python-pptx`, and EML bodies directly from the collapsed thread objects. The parsed text is stored in the pipeline state for downstream use. `storage` uploads each raw file to Google Cloud Storage in parallel, populating the retrieval-by-ID cache used by the `read_attachments` tool.
- `embedding` follows `parsing` and chunks all parsed document text before indexing it into the persistent BigQuery vector store using Google's `gemini-embedding-001` model. This node runs independently and terminates at `END` once indexing is complete.
- `analyze` is a fan-in node that waits for both `parsing` and `initialize_input` to complete before executing. It is the core extraction node: all parsed documents and emails are dispatched to the LLM in size-bounded batches, and the model is asked to extract structured entities — events, claims, damages, deadlines, and document metadata — from each batch using the Pydantic schema described in Section 3.1.6. LLM calls are issued concurrently across batches via `asyncio.gather`, subject to the semaphore concurrency limits. The extracted entities from all batches are then merged into the pipeline state.
- `update_metadata` follows `analyze` and runs a further LLM call to synthesise a coherent case title, background narrative, start date, and primary party list from the full set of extracted data, overwriting the initial seed produced by `initialize_input`. This ensures the factsheet metadata reflects the complete document set rather than only the initial user description.
- `qc_analysis` fans out from `analyze` in parallel with `update_metadata`. It runs a quality-control pass over the extracted entity set, asking the LLM to identify logical inconsistencies, date conflicts, missing required fields, or implausible significance assignments. The results are logged and — if critical issues are found — surfaced to the frontend as warnings, but do not block the pipeline. This node provides a lightweight audit trail of extraction quality without introducing a blocking correction loop.
- `save` follows `update_metadata` and persists all extracted and updated entities to Supabase PostgreSQL via the `SupabaseManager` CRUD interface. Once saved, the factsheet is immediately available for agent queries.

After all terminal nodes (`embedding`, `storage`, `save`, `qc_analysis`) complete, a final status event is streamed to the frontend via SSE confirming that the project is ready for querying.

Update pipeline. The update pipeline shares the same graph topology as initialization, with one structural difference: a `load_project_data` node is inserted at `START` in parallel with `collapse_emails`. This node retrieves the existing factsheet from the database before the new documents are processed. The `analyze` fan-in then waits for both `load_project_data` and `parsing` to complete — substituting the `initialize_input` branch used during initialization — so that the LLM extraction step receives both the new document text and the full prior case state. This allows the model to resolve conflicts with previously extracted entities (e.g., duplicate events or parties already in the factsheet) rather than treating each update as an isolated extraction. The `save` node in the update pipeline merges new entities with the existing factsheet rather than overwriting it.




Figure 3.3: Init and update pipeline graph. `collapse_emails` and `initialize_input` (init) or `load_project_data` (update) run in parallel from `START`. After email extraction and document parsing, `analyze` acts as the central fan-in node. `update_metadata`, `qc_analysis`, `embedding`, and `storage` fan out from there, each terminating independently at `END`.

3.1.8 Clean Pipeline

The `ProjectClean` class provides a cleanup graph that can be triggered on demand (see Figure 3.4):

- **Clean Elements** (`compile_clean_elements`): `load_project` → `clean` → `save_results`. The `clean` node runs `_dedup_elements_node`, which dispatches all element types (events, parties, claims, damages, deadlines) to the LLM in configurable chunks concurrently via `asyncio.gather`. Crucially, the LLM is only asked to return the IDs of elements to keep — rather than rewriting full element objects — which minimises output token usage, reduces hallucination risk, and enables parallel deduplication of several hundred elements efficiently. Elements whose IDs are not returned are discarded, and results are saved back to the database.




Figure 3.4: Clean pipeline graph: `load_project` → `clean` → `save_results`, used for both element deduplication and metadata rewriting.

3.1.9 LLM Configuration

The system supports multiple LLM providers through LangChain’s `init_chat_model` interface. **Google Gemini** serves as the primary provider, with **OpenAI GPT** available as a secondary option. The system also supports a range of open-source models from providers such as Qwen (Alibaba) and Zhipu AI, accessed through compatible APIs. The active model is configurable per session request, allowing for flexible experimentation and evaluation across providers and model families. As the project uses Langchain, new models and providers can be added with minimal code changes if desired for future extensions or experiments.

3.1.10 Tech Stack Summary

Table 3.2 provides a concise overview of all major technologies used in the system, grouped by architectural layer.

Layer	Technology
Language	Python 3.13
Agent framework	LangGraph (StateGraph)
LLM orchestration	LangChain (<code>init_chat_model</code>)
LLM providers	Google Gemini, OpenAI GPT, Together AI
Backend API	FastAPI + Uvicorn
Streaming	Server-Sent Events (SSE)
Frontend	Streamlit
Authentication	Supabase Auth (JWT)
Database	Supabase (PostgreSQL)
Vector store	BigQuery Vector Store (persistent, cross-session)
File storage	Google Cloud Storage (GCS)
Conversation checkpointer	LangGraph <code>AsyncPostgresSaver</code>
Document parsing	Pymupdf4llm, MarkItDown
Embeddings	Google <code>gemini-embedding-001</code>
Web search	Tavily

Table 3.2: System tech stack.

3.2 Design Choices

The system architecture was guided by several pragmatic design choices, prioritizing development speed, flexibility, and cost-efficiency given the constraints of a master’s thesis. By separating the descriptive architecture from the rationale, the core motivations for the selected technologies are outlined below.

3.2.1 Agent Layer and Orchestration

The core agent is implemented using **LangGraph** [9]. This framework was chosen primarily for its fast build time, which is essential given the limited time available for this thesis. It provides a clean **StateGraph** abstraction that is highly expandable for future extensions, such as adding dedicated reasoning nodes or document re-ranking stages. Furthermore, LangGraph is model-agnostic, providing built-in support for multiple LLM providers. This flexibility is crucial, allowing the system to easily switch between models or add new ones as desired for evaluation and experimentation.

3.2.2 Database and Storage Architecture

The system’s persistence layer separates structured relational data, unstructured document storage, and semantic vector retrieval to optimize for specific access patterns and constraints.

Relational Storage (PostgreSQL via Supabase): PostgreSQL was selected for its flexibility and highly efficient reads across tables. The structured nature of the factsheet state requires strict enforcement of relationships through primary and foreign keys to maintain data integrity and orderly matching on IDs [20]. An initial attempt was made using a NoSQL database (Firestore); while easy and fast to set up, the lack of an enforced schema led to a loss of overview, with state inconsistencies and errors accumulating over time. Additionally, performing complex reads and combining data across different entities proved expensive and inefficient in NoSQL.

Supabase was chosen as the PostgreSQL provider because it is affordable, open-source, and seamlessly integrates other necessary products like Authentication. Having these features included out-of-the-box makes the development process straightforward and fast.

Database Schema

The database schema is designed to support the agent’s state management and document metadata, while being flexible enough to accommodate future extensions. There are three main groupings of tables:

Project related tables
Projects
Parties
Events
Deadlines
Attachments
Emails
Claims (legal specific)
Damages (legal specific)

Session related tables
Sessions
Session Events
Session Attachments

Tables for storing document metadata and relationships.

Agent checkpoints
Checkpoints
Writes
Blobs

Tables for the agent’s memory and conversation history.

These tables are combined to form the factsheet state, which is the agent’s internal representation of the case. The factsheet is updated and referenced throughout the agent’s operation, and it is crucial that it is stored in a structured and consistent manner.

File Storage (Google Cloud Storage): Raw documents (PDF, DOCX, EML, etc.) are stored in Google Cloud Storage (GCS). There were no specific architectural requirements for this other than cost-efficiency and reliability. GCS is affordable, integrates cleanly with other Google Cloud services already used in the stack, and provides durable object storage without requiring any dedicated infrastructure.

Relation to Graph-Based Retrieval: The combination of a relational schema and a vector store gives the system a hybrid retrieval architecture that shares structural similarities with graph-based RAG approaches such as GraphRAG [15]. The relational tables — parties, events, claims, damages, and deadlines — function as typed entity nodes, each linked back to the source document through foreign-key relationships and a `source_id` field. These links act as structured edges connecting extracted knowledge to its documentary origin, enabling the agent to reason not just about individual facts but about which documents substantiate them. Semantic search over the raw document chunks (via the vector store) serves the same role as embedding-based traversal in graph retrieval: it handles queries where the relevant content is not captured by the structured entities. The factsheet can therefore be understood as a domain-specific, relational graph over the case — not a general-purpose property graph, but a structured knowledge base with typed entities and cross-document connections, backed by semantic search for unstructured content.

Vector Store and Embeddings: Semantic retrieval relies on **Google’s embedding model** (`gemini-embedding-001`). This model was specifically chosen for its strong multilingual capabilities; smaller and cheaper open-source models often struggle with RAG matching when queries or documents are not strictly in English, which is a critical requirement for handling Norwegian legal texts.

For the vector database, a two-tier approach is used. **ChromaDB** is utilized as an in-memory store for fast, session-scoped retrieval. For persistent, project-scoped retrieval, **BigQuery Vector Search** is employed. BigQuery was chosen because it is affordable and operates on an on-demand pricing model. The primary drawback of BigQuery is its high latency, meaning a significant amount of time is spent indexing and saving to the vector store. However, this trade-off is worthwhile: the alternative would require provisioning an always-running vector database server, which introduces high continuous costs and additional operational setup.

3.2.3 FastAPI and Server-Sent Events

The backend is exposed as a standalone REST API using **FastAPI** [19] to decouple it from the frontend. **Server-Sent Events (SSE)** are used for streaming responses because the agent pipeline is long-running and asynchronous. SSE provides a responsive user experience by pushing incremental status updates and partial results without the complexity of WebSockets, which are unnecessary for this unidirectional communication flow.

3.2.4 Concurrency, Reliability, and Tool Efficiency

Several engineering decisions were made to ensure the system operates efficiently under the latency and throughput constraints imposed by external APIs and large document volumes.

Async-first architecture. The entire backend stack is built on Python’s `asyncio` event loop. FastAPI natively supports async request handlers, and all LangGraph pipeline nodes are declared as `async` coroutines. This allows the server to handle multiple concurrent sessions without blocking on I/O-bound operations such as LLM calls, database queries, or storage reads.

Semaphore-based concurrency control. To prevent overloading external services, four independent `asyncio.Semaphore` instances are used — one per resource class: LLM, database, storage, and vector store. Each semaphore has a configurable maximum concurrency value drawn from `config.toml` (e.g. `[async.llm] max_concurrent_requests = 20`). Pipeline nodes acquire the relevant semaphore before dispatching work, ensuring that bursts of parallel tasks are automatically throttled without busy-waiting or sleep loops.

Batch-wise and chunk-wise parallel processing. During project initialization, documents are grouped into batches before being dispatched for LLM analysis. Batching is driven by two configurable thresholds: a maximum byte-size per batch (`project.threshold`) and a maximum file count per batch (`project.max_attachments` / `project.max_emails`). Each batch is wrapped in an async task that acquires the LLM semaphore, and all tasks are dispatched concurrently via `asyncio.gather()`. The same chunk-wise strategy is applied during factsheet cleaning and deduplication: element lists (events, claims, damages, etc.) are split into fixed-size chunks (`project.clean.chunk_size`), each chunk is cleaned by a separate concurrent LLM call, and results are merged after all tasks complete. This approach keeps individual prompts within token limits while maximising throughput.

Rate limiting and 429 handling. To comply with provider-imposed rate limits and avoid HTTP 429 (Too Many Requests) errors, each LLM instance is wrapped with LangChain’s `InMemoryRateLimiter`. The limiter implements a token-bucket algorithm: tokens are refilled at a configurable rate (`async.llm.requests_per_second`) and a burst capacity (`async.llm.max_burst_size`) allows short spikes. A single stateful limiter instance is shared across all coroutines that use the same model, so the rate limit is enforced globally rather than per-request.

Retry mechanism with exponential backoff. Transient API errors (e.g. network timeouts or momentary 5xx responses) are handled by a dedicated `apply_retry()` helper. It wraps any LangChain chain with `with_retry()`, configuring a stop condition (`retry_attempts` from config) and exponential backoff with random jitter (`wait_exponential_jitter=True`). The retry decorator is applied *after* structured-output binding so that validation failures are also covered. If all retry attempts are exhausted, the error is propagated and logged.

TOML-based configuration management. All tunable parameters — concurrency limits, rate limiter settings, retry counts, chunk sizes, model endpoints, and storage bucket names — are centralised in a single `config.toml` file. On startup, the file is parsed into a fully typed `AppConfig` object built with `Pydantic` and `pydantic-settings`. Nested configuration sections map directly to typed dataclasses (e.g. `AsyncConfig`, `ProjectConfig`, `VectorstoreConfig`), providing IDE auto-complete, validation at load time, and a single source of truth for the entire application. A `get_app_config()` singleton cached with `@lru_cache` ensures the config is parsed only once per process.

Tool interface design — string IDs over full objects. Agent tools are designed to accept minimal, flat inputs: all tools receive only string identifiers (e.g. `project_id: str, ids: list[str]`) rather than full domain objects. This keeps the JSON schema presented to the LLM small, reduces the token footprint of tool calls, and makes it straightforward for the model to construct valid invocations. Internal data fetching (loading full attachment records, querying the database) is performed inside the tool implementation, hidden from the model entirely.

Context compression via filtering and display shortening. A key design concern for the custom agent is managing the size of the structured context injected into each LLM call. As a case accumulates material, a naive approach of injecting the full factsheet verbatim would cause token consumption to grow unboundedly. Two complementary compression mechanisms are applied to mitigate this.

First, the `show_elements` and `list_attachments` tools accept optional `significance` and `date_range` parameters. These filters are implemented as PostgreSQL RPC functions on the Supabase backend and are applied server-side before any data is returned to the agent, ensuring that only elements above a configurable significance threshold and within the relevant time window are included in the context.

Second, the string representation of each structured element is deliberately abbreviated. Rather than serialising the full object with all fields, a compact display format omits less informative fields and truncates long free-text descriptions. Together, these two mechanisms reduce the effective model input by approximately 70–80% compared to an unfiltered, full-fidelity serialisation of the same factsheet. This compression is a significant contributor to the finding that the custom agent’s average token consumption remains statistically indistinguishable from that of the simpler baseline agent despite carrying substantially richer case context.

3.3 Evaluation Design

3.3.1 Datasets

The evaluation relies on two datasets developed specifically for this thesis in collaboration with legal experts from the law firm BAHR AS. Each dataset represents a real legal case processed through the Norwegian court system.

THRD-2021-163881 is based on a publicly available district court case concerning a residential property purchase in which significant latent defects were discovered after handover. Shortly after taking possession, a recurring water leak in the rental unit resurfaced — one that had been reported and allegedly repaired prior to sale. Subsequent inspections revealed structural deviations in the concrete slab and that the rental annex partly encroached on the neighbouring property. The buyers formally notified the seller of their intent to rescind and subsequently filed a writ of summons against both the seller and the property’s insurer. A mediated settlement was reached but later rescinded after a new expert report concluded that the concrete slab required full demolition and reconstruction. The case was reinstated and proceeded to a main hearing with expert testimony on the slab, property boundaries, and unapproved structural elements, resulting in a final judgment.

The authentic court record — comprising the final party submissions and the court’s verdict — is used directly as part of the document set. The remaining case materials (correspondence, internal notes, technical reports, agreements, and supporting evidence) were not publicly available and were therefore synthetically reconstructed to match the factual matrix and legal reasoning documented in those authentic records. The synthesis was performed using Claude.ai’s consumer interface rather than the professional API; unlike API access, the consumer product may use submitted content to train future models. To mitigate privacy risks, all personal names and identifying particulars in the synthetic documents were replaced with randomised fictional identities before any material was submitted for synthesis. This hybrid approach provides a realistic document topology while respecting both confidentiality and data-protection considerations. However, synthetic documents carry inherent limitations: they may not fully replicate the linguistic register, internal inconsistencies, and informational density characteristic of authentic legal correspondence and technical reports. Retrieval signals may consequently be somewhat cleaner than in a fully authentic case, which could modestly inflate measured performance relative to production conditions. This limitation is discussed further in Section 3.3.2.

TOSL-2024-125319 is a fully authentic case concerning a personal injury sustained at a commercial ski resort’s snowpark. The claimant suffered a severe traumatic brain injury following a jump accident, resulting in extended hospitalization, intensive rehabilitation, and a subsequent finding of full and permanent disability. The claim is brought against the resort’s liability insurer, who contests liability on the grounds that the facility met applicable industry standards and that no prior accidents had been recorded at the specific jump. A central evidentiary dispute concerns the design and structural specification of the jump at the time of the accident, and its subsequent removal and reconstruction shortly thereafter. The case involves multiple competing expert opinions on industry safety standards and causation.

Explicit permission was obtained to use the original case documents. The dataset is substantially larger and more complex than THRD-2021-163881: it involves more parties, a longer contested timeline, more contested legal issues, and a higher total document volume. The original case file includes close to a thousand emails; these were excluded from the evaluation dataset due to the prohibitive per-run cost and execution time they would introduce across repeated evaluation trials. Even after this exclusion, TOSL-2024-125319 remains considerably larger and more complex than THRD-2021-163881. Because it consists entirely of authentic documents, it more closely represents the conditions under which the agent would operate in production, and its results should be given greater interpretive weight.

The datasets are organized to simulate the chronological lifecycle of a real use-case example. Each case is structured into several *sessions*, each containing a sequence of *queries*. This structure is designed to evaluate two critical capabilities of the agent:

1. **Contextual Awareness:** The ability to maintain a coherent understanding of the case as it evolves through different stages of litigation.
2. **Analytical Accuracy:** Each query is paired with a "ground truth" expected output, defined in collaboration with legal experts. These references serve as the ground truth for measuring

correctness, relevancy, and completeness.

By following a chronological order, the sessions require the agent to handle increasing amounts of information and maintain consistency over time, reflecting the professional requirements of a legal practitioner managing an ongoing case.

Constructing these datasets required substantial effort, as no publicly available alternative combined the necessary structural elements: session-based chronological queries grounded in large, multi-document case files with expert-verified ground-truth answers. A survey of existing legal and large-scale QA datasets confirmed this gap: while many share some characteristics with the present work, none combine all the necessary elements. For instance, datasets like **Multi-LexSum** [21] offer large case files and expert summaries but lack the session-based, chronological query structure needed to evaluate temporal reasoning. Others, such as **CaseGen** [22], provide authentic, multi-section case samples but do not include the large, unstructured document collections typical of real-world litigation. Datasets like **EnronQA** [23] contain extensive email corpora but are not situated within a legal context. Finally, established benchmarks like **LegalBench** [24], **LexGLUE** [25], and others are primarily focused on legal reasoning over statutes or case law rather than fact management in an evolving case file. The evaluation thus necessitated a custom-built dataset format that included not only a large corpus of interconnected documents and emails but also a series of chronologically ordered, session-based queries designed to test contextual awareness over time.

The construction of the queries and their corresponding ground-truth answers was a meticulous process. Queries were designed to be fact-based, verifiable against the source documents, and challenging enough to require cross-document synthesis. Crafting the ground-truth answers proved substantially more difficult than anticipated.

The initial approach was to write comprehensive reference answers covering all relevant information a legal practitioner might include. In practice, for a large case with hundreds of documents, almost any additional fact can be incorporated without being wrong — the amount of information that is *technically correct to include* is nearly unbounded. Trial evaluation runs revealed that this made automated grading unreliable: DeepEval would mark a model response as incomplete not because it was wrong, but because it omitted details present in an overly comprehensive reference that no practitioner would reasonably demand in a single answer. This introduced systematic noise into the completeness and correctness scores.

The framework was revised iteratively in response to these observations. The goal became to define a *minimal, essential* answer — one that captures the core required facts, ensuring that the evaluation penalises omissions of genuinely important information rather than differences in coverage depth. The factual versus summary query split (see below) was similarly introduced after observing that automated judges reliably penalised structurally different but factually correct responses to open-ended questions. Both adaptations were not part of the original evaluation design but emerged as necessary corrections during the evaluation process; they should be understood as pragmatic responses to limitations in the automated evaluation framework rather than ideal design choices.

This required an in-depth manual review of all case documents for each query. After cleaning and removing anomalies, a total of 742 query-answer pairs were collected for the evaluation. While 742 pairs is not a trivially small sample size, the dataset draws from only two distinct cases, which limits the generalisability of the findings. The queries reflect the factual structure and legal complexity of two specific Norwegian civil disputes; whether the observed performance patterns hold across different jurisdictions, case types, or document styles remains an open question. Expanding the evaluation to cover a broader set of cases is among the most important directions for future work (see Section 6.4).

Query type taxonomy: factual vs. summary. Each query in the dataset is classified into one of two types: *factual* or *summary*. This distinction drives which metrics are computed for each query and is central to the validity of the evaluation.

Factual queries ask for specific, enumerable facts whose correctness can be objectively verified against the case documents: for example, “Who are the parties to the main contract?”, “What damages are claimed by the plaintiff?”, or “On what date was the court submission filed?”. The defining characteristic is not that the answer comes from a single document — many factual queries still require synthesising information across multiple sources — but that the answer is bounded and its required content is well-defined. There is a finite set of facts that a complete and accurate response must cover,

and that set can be stated as a ground-truth reference without imposing a framing choice. Reference-based metrics (correctness, completeness) are reliable on factual queries for this reason: the reference answer is an objective enumeration of required facts, not merely one of many valid formulations.

Summary queries ask the agent to synthesise a narrative overview: for example, “Summarise the key events of the dispute”, or “What is the overall legal position of the defendant?”. These queries do not have a single correct answer. A practitioner could frame an equally valid and complete response in multiple ways — emphasising different events, using different levels of abstraction, or organising the information differently — and all such responses would be professionally appropriate. Reference-based metrics are *not* reliable on summary queries because they compare the agent’s response against one specific reference framing. A response that is factually complete and legally accurate but structured differently from the reference will be penalised as incomplete or incorrect, not because it is wrong, but because it chose a different framing. This constitutes unfair penalisation and would systematically underestimate both agents’ performance on summary questions.

The practical consequence for the evaluation is that completeness is only computed on factual queries, where the concept of “required facts” is well-defined and the metric is interpretable. Correctness and relevancy are computed across all query types, since these metrics are more robust to framing variation: a factually accurate response tends to score well on correctness even if its structure differs from the reference, and relevancy only requires that the response is on-topic. This design decision deliberately sacrifices completeness coverage for measurement validity.

Examples The following examples illustrate the distinction between factual and summary queries and the form of the corresponding ground-truth answers.

Summary query — THRD-2021-163881, Session 1

Query: Provide a short and concise summary of the factual background and development of the case so far, based on the documents uploaded. Focus on the most important events and issues. What is the core of the case?

Reference answer:

- The buyers purchased the property with transfer of ownership on ...
- The seller reported a leak in the rental unit ...
-

Factual query — THRD-2021-163881, Session 1

Query: What information has the buyer received and obtained regarding the leak, both in writing and orally, up to today?

Reference answer:

- ...

Factual query — TOSL-2024-125319, Session 3

Query: Present the experts’ assessments of the jump and compare them with the parties’ arguments.

Reference answer:

- ...

Summary query — TOSL-2024-125319, Session 3

Query: Describe the claimant’s injuries, the progression of the condition, and their state of health from the time of the accident up to the date of filing the claim. Has the condition improved, stabilized, or worsened?

Reference answer: ...

3.3.2 Evaluation Limitations and Caveats

Several factors limit the generalisability of the evaluation and should be considered when interpreting the results.

Absence of a plain LLM baseline. A natural comparison point would be an ordinary LLM queried without any specialised architecture — no document processing, no vector store, no RAG, simply the raw case files passed as context. This approach was explicitly attempted. In practice, the aggregated size of a realistic legal case file — spanning contracts, correspondence, technical reports, and legal submissions — immediately exhausted the rate limits of every model tested, well before a meaningful response could be produced. The sheer volume of document content triggers the providers’ token-per-minute limits even for a single query, making repeated evaluation across sessions entirely infeasible. Feeding entire case files into a plain LLM context is therefore not a viable strategy for any of the models evaluated, and the minimum meaningful comparison is a RAG-enabled architecture. This is also why the evaluation compares the custom agent against a Baseline + RAG configuration rather than an unaugmented LLM.

Dataset size and the scalability trade-off. The evaluation uses a constrained number of cases and queries, as running a large-scale benchmark across multiple model-agent pairs with repeated trials becomes prohibitively expensive and time-consuming. This is, however, in tension with a fundamental property of the system: the agent is designed to improve with larger and richer document sets, as more material increases the factsheet’s coverage. Evaluating on small datasets therefore likely *underestimates* the performance advantage of the custom agent in production-scale deployments.

Heterogeneity between cases and synthetic document limitations. The two evaluation datasets differ substantially in origin, complexity, and document volume. THRD-2021-163881 uses a hybrid document set (authentic court record supplemented by synthetically reconstructed case materials), while TOSL-2024-125319 consists entirely of authentic documents. TOSL-2024-125319 is also considerably more complex: it covers a larger total document volume.

The synthetic component of THRD-2021-163881 is a meaningful limitation. Synthetically generated correspondence and reports are internally consistent and stylistically uniform in ways that authentic legal documents are not: they lack the ambiguity, redundancy, contradictions, and variable information density that accumulate naturally over the lifecycle of a real case. This may make retrieval and extraction easier on THRD-2021-163881 than on a fully authentic case, potentially inflating both agents’ scores and compressing the performance gap between them. Results on THRD-2021-163881 should therefore be treated as indicative rather than definitive. TOSL-2024-125319, consisting entirely of real documents, is the more ecologically valid test of the architecture’s capabilities, and its results carry greater interpretive weight.

Exclusion of pipeline costs from resource metrics. The resource metrics reported in this evaluation — token usage and inference latency — cover only query-time activity: the LLM calls made by the agent when responding to a user query. The one-time cost of running the initialization and update pipeline (document parsing, entity extraction, embedding, and storage) is deliberately excluded.

This exclusion is methodologically correct for two reasons. First, including pipeline costs in a per-query resource comparison would produce a misleading figure on the evaluation datasets. Each dataset contains a small number of evaluation queries; amortising the full extraction cost across, say, four queries makes the cost per query appear orders of magnitude higher than it would be in practice. A production deployment handles hundreds or thousands of queries per case, over which the fixed extraction cost becomes negligible. There is no principled way to choose an amortisation denominator within the scope of this evaluation, and any choice would be arbitrary.

Second, the pipeline does not run at query time. It is an offline preprocessing step that can be triggered asynchronously — overnight, during off-hours, or on document arrival via integrations with document management systems and email. The practitioner does not wait for it. The relevant cost comparison is therefore query-time cost, which is what the reported metrics reflect.

Limitations of expected-answer evaluation. Model responses are graded against expected answers from the QA-pairs produced in collaboration with legal experts at BAHR AS. This approach carries two inherent risks: the expected answers themselves may contain inaccuracies, and — more fundamentally — many legal queries are open-ended, such that a correct and complete model response may differ substantially from the reference without being wrong. The framing problem is particularly acute for summary queries: a response that synthesises the case accurately but organises the information differently from the reference will appear incomplete to an automated judge even though it may be professionally superior. Four mitigations were applied to reduce this risk:

1. **Query type separation:** The dataset distinguishes factual from summary queries (see above). Completeness evaluation is restricted to factual queries, where a single reference answer is an objective ground truth. Summary queries are included in correctness and relevancy evaluation but not in completeness evaluation, preventing framing-based penalisation from inflating the apparent performance gap.
2. **Fact-based scope:** Even within factual queries, the evaluation focuses on document-grounded questions (e.g. “Who are the parties?”, “What damages are claimed?”), deliberately excluding normative or predictive legal questions where correct answers are inherently ambiguous.
3. **Validatable content:** Questions are formulated so that their answers can be directly verified against the case documents, reducing the space for legitimate variation.
4. **Minimal-requirement framing:** Expected answers are written as *minimum* acceptable responses rather than exhaustive ideal answers, so a model response that covers the required facts is scored as correct even if it includes additional relevant information.

3.3.3 Evaluation Environment

Evaluation Framework:

The evaluation will be conducted using **DeepEval** [26], a framework for evaluating LLM outputs against reference answers. DeepEval allows for automatic scoring based.

Agent Configurations:

Two agent configurations are evaluated:

- **Custom Agent:** The full agent configuration with access to all tools, including structured factsheet access (`show_elements`, `list_attachments`), direct file reading (`read_attachments`), and RAG capabilities.
- **Baseline + RAG Agent:** An ablated version with access to web search, law retrieval, and project document RAG, but without direct file access or structured factsheet tools.

A plain LLM baseline without document retrieval was considered but excluded from the evaluation for the reasons described in Section 3.3.2: the volume of case material immediately saturates provider rate limits, making repeated evaluation impossible.

Model Configuration:

Since the scope of the thesis is not to evaluate the models themselves, but rather the design and architecture of the agent, a selection of four models is used for comparison. These include both closed-source and open-source variants to ensure the agent’s performance is consistent across different provider architectures.

For open-source models, several were tested. Initial trials with smaller models, such as **Qwen3-Next-80B-A3B-Instruct** [27], demonstrated that they lacked the capacity to handle the complexity of the reasoning and extraction tasks, failing to produce coherent or useful outputs. This became especially evident for the custom model, that heavily relies on the quality of the extracted entities. As a consequence, these smaller models were excluded from the evaluation. Consequently, the evaluation focused on two of the top-performing large-scale open-source models available at the time of writing: **Qwen 3.5** from Alibaba and **GLM-5** from Zhipu AI. Both have a similar scale in terms of total parameters (approx. 400B), but differ significantly in their Mixture-of-Experts (MoE) architecture. Qwen activates 17B parameters

per forward pass, whereas GLM-5 activates 40B, giving it a theoretical advantage in capacity that is later reflected in the results. Inference for these models was run via the `together.ai` [28] API, which was chosen for its production-grade scalability, affordable pricing, and seamless compatibility with the existing LangChain and OpenAI SDKs.

For closed-source models, Google’s Gemini family was selected. **Gemini 2.5 Flash** provides a cost-effective and reliable baseline, suitable for extensive testing without incurring prohibitive costs. **Gemini 2.5 Pro** serves as the high-capability "gold standard". While other providers were considered and are supported by the system’s architecture, Google’s offerings were particularly advantageous for this evaluation due to their generous rate limits (per second, minute, and day) and large context windows. This was a critical practical consideration, as it minimised the risk of hitting API limits during the time-consuming data collection process, which could otherwise halt evaluation runs despite the presence of retry logic.

Table 3.3: Overview of Large Language Models used in evaluation.

Model Name	Type	Total Params	Active Params	Provider
Gemini 2.5 Pro [29]	Closed	Unknown	Unknown	Google
Gemini 2.5 Flash [?]	Closed	Unknown	Unknown	Google
Qwen 3.5 [30]	Open (MoE)	397B	17B	Alibaba
GLM-5 [31]	Open (MoE)	435.2B	40B	Zhipu AI

Table 3.3 lists the four models included in the evaluation, together with their parameter counts and providers.

The evaluation will be conducted across all datasets using consistent model-agent pairings. To account for the inherent stochastic nature of LLM outputs, each evaluation set (model + agent configuration) will be executed 3 times to ensure statistical significance in the results.

3.3.4 Metrics

The evaluation focuses on two categories of metrics: quality metrics that assess the accuracy and completeness of the agent’s responses, and resource metrics that assess the computational cost of producing those responses. Together they allow the main research question to be evaluated jointly — whether any quality gains are proportionally offset by increased resource consumption.

Correctness, Relevancy, and Completeness (DeepEval). These three semantic quality metrics are evaluated using the **DeepEval** framework [26], which uses a separate LLM as an automated judge to score agent outputs against gold-standard reference answers. The judge model used throughout this evaluation is **Gemini 2.5 Flash**: it is cost-efficient enough to run at the scale required (742 query-answer pairs $\times$ 3 runs $\times$ 3 metrics), well-documented and widely used. Each response is evaluated independently on a continuous scale from 0 to 1.

Correctness measures whether the factual claims in the agent’s response are accurate relative to the reference answer. It is the most diagnostic metric in this evaluation because it directly captures the architectural advantage under test: a structured factsheet provides the agent with pre-extracted, relational facts that vector retrieval alone may fail to surface from a large document corpus.

Relevancy measures whether the response addresses the question asked, rather than drifting to adjacent or tangential content. It functions primarily as a sanity check: as long as an agent engages substantively with the question, the response tends to be on-topic regardless of the underlying architecture. Cases where relevancy fails are therefore informative precisely because they indicate a retrieval or engagement failure, not merely a factual gap.

Completeness measures whether the response covers all facts required by the reference answer. It complements correctness by penalising omissions rather than factual errors: a response can be entirely accurate in what it states while still omitting important facts that the reference includes. Together, the three metrics span the accuracy–coverage–focus dimensions of answer quality.

A response is classified as a *pass* if its aggregate score meets or exceeds a predefined quality threshold; otherwise it is a *fail*. The proportion of queries receiving a passing score is reported as the **Passrate** metric. The passrate threshold used in the evaluation pipeline is 0.5.

BERTScore and Sentence-BERT Similarity. BERTScore [32] quantifies token-level semantic overlap between the agent’s response and the reference answer using contextualised BERT embeddings. It computes *precision* (the fraction of generated tokens that are semantically present in the reference), *recall* (the fraction of reference tokens covered by the generated response), and an *F1* score that balances the two. Unlike the LLM-judge metrics, BERTScore requires no separate inference call and is insensitive to judge-prompt phrasing or judge model choice, making it a useful cross-validation signal that is independent of judge-model behaviour.

Sentence-BERT similarity complements BERTScore by capturing global embedding-level proximity between the two texts as a whole, rather than aligning individual tokens. The distinction matters because a response can score well on BERTScore — meaning its surface vocabulary overlaps with the reference — while omitting structured factual content that correctness and completeness evaluation would penalise. Including both metrics provides a more complete picture of semantic alignment and guards against the possibility that quality differences visible in the LLM-judge metrics are an artefact of embedding granularity.

Self-consistency. Self-consistency measures how semantically similar the agent’s responses are when the same query is posed in different sessions. Because each agent run involves some inherent stochasticity (temperature sampling, retrieval variance, session-level context differences), the same query may receive different responses across runs even if the underlying case state is identical. A high self-consistency score indicates that the agent converges reliably on the same answer regardless of session context; a low score suggests sensitivity to retrieval noise or prompt variation. This metric captures a dimension of reliability that per-query quality scores cannot: not just whether the agent answers correctly, but whether it does so predictably.

Self-consistency is computed as the mean pairwise cosine similarity of sentence embeddings across all runs of the same query:

$$\text{Self-Consistency} = \frac{2}{n(n-1)} \sum_{i=1}^{n-1} \sum_{j=i+1}^n \cos(\mathbf{e}_i, \mathbf{e}_j) \quad (3.1)$$

where n is the number of runs, $\mathbf{e}_i$ and $\mathbf{e}_j$ are the sentence embeddings of the i -th and j -th responses, and cosine similarity is defined as:

$$\cos(\mathbf{e}_i, \mathbf{e}_j) = \frac{\mathbf{e}_i \cdot \mathbf{e}_j}{\|\mathbf{e}_i\| \|\mathbf{e}_j\|} \quad (3.2)$$

Resource Metrics. Four resource metrics are recorded for each agent run: **inference duration** (wall-clock seconds from query submission to final response), **input token count**, **output token count**, and **LLM call count**. Input tokens directly drive API cost and context-window pressure; inference duration determines user-facing latency; LLM call count indicates how much internal orchestration each query requires. These metrics are analysed alongside quality scores to determine whether the custom agent’s architectural overhead — the structured factsheet context injected on each call — imposes a measurable additional cost relative to the simpler baseline.

Additionally, an estimated cost per run is computed using each provider’s advertised per-token prices:

$$\text{cost} = \text{input_tokens} \times p_{\text{in}} + \text{output_tokens} \times p_{\text{out}}$$

where p_{in} and p_{out} are the input and output token prices respectively. Because absolute costs differ substantially between datasets (TOSL-2024-125319 involves larger document volumes and longer sessions than THRD-2021-163881), raw cost figures are normalised to within-dataset percentile ranks for cross-model comparison.

3.3.5 Research Question Operationalization

Each research question is operationalized through a specific subset of the evaluation evidence collected.

Main RQ — accuracy improvement without proportional cost increase. The main RQ is evaluated by jointly examining quality metrics (correctness, relevancy, completeness, passrate) and resource metrics (input tokens, inference duration) across all queries and sessions. A positive finding

requires both a statistically significant quality advantage for the custom agent *and* the absence of a significant resource cost increase. The one-sided Mann-Whitney U tests described below provide the formal statistical basis.

RQ1 — mitigation of cross-document retrieval limitations. RQ1 is addressed by examining the quality metrics specifically on queries that require synthesis across multiple documents and sessions. The dataset-level breakdown (Section 4) isolates TOSL-2024-125319, which contains a larger document volume and more contested cross-document relationships, as the primary test case for cross-document reasoning. A larger performance gap on TOSL-2024-125319 compared to THRD-2021-163881 is interpreted as evidence that structured extraction specifically addresses the multi-document limitation.

RQ2 — performance stability across growing document volume. RQ2 is addressed by inspecting session-level performance trajectories. The datasets are organized chronologically into multiple sessions, each adding new documents. A custom agent that maintains or improves quality across sessions while the baseline degrades or stagnates provides evidence that the structured representation scales gracefully with document volume.

RQ3 — model capacity compensation. RQ3 is addressed by examining performance per model. The evaluation includes models of varying capability, from smaller open-source models (Qwen3 MoE 3B active parameters) to larger closed-source models (Gemini 2.5 Flash, Claude Sonnet). If the custom agent enables a smaller model to reach performance levels comparable to a larger baseline model, this constitutes evidence that structured context can partially substitute for raw model capacity.

3.3.6 Normality Testing

Before selecting hypothesis tests, the distributional properties of the metric scores were assessed using the **Shapiro-Wilk test** for normality. The Shapiro-Wilk test evaluates whether a sample could plausibly have been drawn from a normal distribution; a significant result ($p < 0.05$) rejects normality. Applied to each metric across both agent configurations, the test returned $p < 10^{-18}$ in every case, providing overwhelming evidence against normality. Visual inspection confirmed the cause: all metrics exhibit pronounced left-skew with heavy concentration near 1.0 and a long lower tail (see Figure 3.5). This rules out parametric tests that assume normally distributed residuals.

3.3.7 One-sided Hypothesis Tests

Given the non-normal distributions established above, all hypothesis tests are conducted using the one-sided **Mann-Whitney U test**, a non-parametric rank-based test that assesses whether one distribution tends to produce stochastically larger values than another, without assuming normality.

For quality/performance metrics (correctness, relevancy, completeness, passrate, self-consistency, BERTScore):

$$H_0 : \eta_{\text{custom}} \leq \eta_{\text{baseline+RAG}} \quad H_1 : \eta_{\text{custom}} > \eta_{\text{baseline+RAG}}$$

For cost/efficiency metrics (token usage, inference time):

$$H_0 : \eta_{\text{custom}} \geq \eta_{\text{baseline+RAG}} \quad H_1 : \eta_{\text{custom}} < \eta_{\text{baseline+RAG}}$$

Notation:

- η_{custom} = population median for the custom agent
- $\eta_{\text{baseline+RAG}}$ = population median for the baseline with RAG agent

All tests are one-tailed / one-sided.




Figure 3.5: Distribution of correctness scores across all queries and sessions. The pronounced left-skew and heavy concentration near 1.0 violate the normality assumption, motivating the use of non-parametric Mann-Whitney U tests.

3.3.8 Levene’s Test for Equality of Variance

The Mann-Whitney U test evaluates whether one distribution tends to produce stochastically *larger* values than another, but says nothing about whether the two distributions differ in *spread*. To test whether the custom and baseline agents exhibit significantly different score variability, Levene’s test for equality of variance [33] is applied to each quality metric.

The test statistic is:

$$W = \frac{N - k}{k - 1} \cdot \frac{\sum_{i=1}^k N_i (\bar{Z}_{i.} - \bar{Z}_{..})^2}{\sum_{i=1}^k \sum_{j=1}^{N_i} (Z_{ij} - \bar{Z}_{i.})^2} \quad (3.3)$$

where:

- k = number of groups (here $k = 2$: custom and baseline)
- N = total number of observations across both groups
- N_i = number of observations in group i
- $Z_{ij} = |Y_{ij} - \bar{Y}_{i.}|$, the absolute deviation of observation j in group i from its group mean
- $\bar{Z}_{i.}$ = mean of Z_{ij} within group i
- $\bar{Z}_{..}$ = overall mean of all Z_{ij}

Under the null hypothesis of equal variances, W follows an F -distribution with $k - 1$ and $N - k$ degrees of freedom. A significant result ($p < 0.05$) rejects the null and indicates that the two groups differ significantly in spread. The hypotheses are:

$$H_0 : \sigma_{\text{custom}}^2 = \sigma_{\text{baseline+RAG}}^2 \quad H_1 : \sigma_{\text{custom}}^2 \neq \sigma_{\text{baseline+RAG}}^2$$

The standard parametric test for equality of variances is **Bartlett’s test**, which evaluates whether two or more groups share the same population variance under the assumption that each group is drawn from a normal distribution. Because the metric distributions are non-normal (established above),

Bartlett’s test is not applicable here. **Levene’s test** [33] is therefore used instead: it tests the same null hypothesis — equal variances across groups — but replaces the normality assumption with a median-based absolute-deviation statistic, making it robust to the skewed distributions observed in the evaluation data.

3.4 Development Tooling and AI Assistance

This section documents the tools and practices used during the development of this thesis — both the software infrastructure supporting the project and the role of AI assistance in the process.

3.4.1 Version Control and Repository Structure

The entire project is maintained in a single Git monorepo hosted on GitHub. This encompasses the agent source code, auxiliary modules (`ui`, `evals`, `data-viewer`), the evaluation datasets, and the $\text{\LaTeX}$ source for this thesis. Keeping all artefacts in one repository simplifies cross-referencing between the written thesis and the underlying implementation, and allows figures generated in the evaluation notebooks to be committed directly to the `thesis/figures/` directory — making figure updates a matter of running a notebook cell and committing the result. The only assets not stored in the repository are the datasets and case files, which are hosted in Google Cloud Storage due to their size and confidentiality.

Git’s branching model and worktree support were used throughout development. Feature branches isolated experimental work from the stable main branch, and Git worktrees allowed parallel development — for example, running an evaluation trial on one worktree while implementing new features in another — without disrupting the working state.

3.4.2 Dependency Management

Python dependencies are managed with `uv`, a fast, standards-compliant package manager. `uv`’s workspace feature allows all four Python packages (`agent`, `ui`, `evals`, `data-viewer`) to share a single lock file and virtual environment root while maintaining separate dependency declarations per package. This was a key practical consideration for a monorepo of this scope.

3.4.3 Report Authoring

The thesis is written in $\text{\LaTeX}$ and compiled using `latexmk`, which automates the multi-pass build required for references, cross-references, and bibliography. The source is version-controlled alongside the codebase, enabling tight coupling between the report and the implementation: code snippets, figures, and table data are all traceable to specific commits.

3.4.4 Data Viewer

A lightweight **Streamlit** application (`data-viewer`) was developed to support dataset work and result inspection. It provides a browser-based interface for browsing and editing both training datasets and case files stored in GCS, displaying evaluation results, and inspecting analysis outputs. This reduced the overhead of dataset quality control and made it easier to identify issues during ground-truth annotation.

3.4.5 Use of AI Assistance

AI tools were used during the development of this thesis in accordance with NMBU’s guidelines for the use of artificial intelligence in academic work [34]. All use is disclosed here as required.

The primary tool was **Claude Code** (Anthropic), a CLI-based AI assistant integrated into the IDE.

Code. Claude Code was used for code review, targeted refactoring, and implementation of cross-cutting infrastructure concerns — such as retry logic, rate-limit handling, and concurrency patterns — that are peripheral to the core research contribution. All architectural decisions, data model design, and system-level trade-offs were made independently by the author.

Report writing. AI assistance was used for reformulation and rewording of written content, improving clarity and maintaining a consistent academic register. It was also used to locate relevant background literature and identify applicable theoretical frameworks. All substantive content, arguments, and conclusions are the author’s own, and all AI-assisted formulations have been critically reviewed and verified.

Dataset construction. For the THRD-2021-163881 dataset, synthetic case materials (correspondence, internal notes, and supporting documents) were generated using an LLM, seeded from the authentic court record. This approach is discussed in Section 3.3.2 as a methodological caveat: while the synthetic documents were constructed to be consistent with the factual matrix of the authentic record, there is an inherent risk of logical inconsistencies between the real and synthetic materials, and this constitutes a recognised limitation of that dataset.

Ground-truth answers were reviewed and validated using AI-assisted quality control, in collaboration with legal experts at BAHK AS.

Chapter 4

Results

Data pre-processing. Prior to analysis, runs with a duration exceeding 600 seconds (10 minutes) were excluded as anomalies. These correspond to timed-out or otherwise pathological executions whose resource measurements would distort the distribution of duration and token statistics. All results reported below are computed on the filtered dataset.

4.1 Quality Metrics

4.1.1 DeepEval: Correctness, Relevancy, and Completeness

The joint distribution of correctness, relevancy, and completeness (defined in Section 3.3.4) across all queries and sessions is shown in Figure 4.1.




Figure 4.1: Pairplot of DeepEval quality metrics (correctness, relevancy, completeness). Diagonal panels show per-metric density estimates. The custom agent (orange) exhibits higher correctness and completeness, while relevancy is similarly high for both configurations.

Both distributions are left-skewed across all three metrics, with density concentrated near 1.0 — but the degree of concentration differs substantially. The custom agent’s distributions are more tightly clustered near the upper end, while the baseline’s are more spread, with a heavier low-scoring tail. This confirms that the custom agent’s advantage is not driven by a small number of high-scoring queries but reflects a broad shift in the distribution. Relevancy scores are highly concentrated near 1.0 for both configurations — as long as an agent attempts to answer the question, the response tends to be on-topic regardless of architecture — though the baseline’s low-scoring tail corresponds primarily to cases where retrieval fails entirely. Completeness follows the same pattern: both are left-skewed, but the custom agent’s distribution is more concentrated near 1.0, while the baseline exhibits more low-completeness responses. Overall, the pairplot reveals greater consistency and reliability in the custom agent’s responses across all three quality dimensions.

Session-level trajectories tell the same story longitudinally (Figure 4.2).




Figure 4.2: DeepEval quality metrics across sessions for each agent configuration.

Both agents broadly track each other through the session progression, but the custom agent remains consistently above the baseline throughout. The shaded bands represent 95% confidence intervals across the three repeated runs; the narrower bands on the custom agent’s trajectory reflect a tighter spread across runs, consistent with the lower score variance reported in the statistical analysis. This session-level consistency reinforces the pairplot finding that the quality advantage is broad and reliable rather than driven by occasional strong queries.

Breaking the correctness trajectory down by dataset (Figure 4.3) reveals that the performance gap is not uniform across cases.




Figure 4.3: Correctness over session progression, split by dataset.

On TOSL-2024-125319 — the fully authentic case with the larger document volume and greater factual complexity — the custom agent’s advantage grows more noticeably as sessions accumulate. On THRD-2021-163881, the custom agent also leads, but the gap is smaller and more stable. This pattern is consistent with the architectural hypothesis: the factsheet and structured retrieval tools confer the greatest benefit precisely when the case is large and complex, where naive RAG retrieval is most prone to missing relevant facts dispersed across many documents.

Figure 4.4 shows the same comparison as a violin plot, making the distributional differences and interquartile ranges directly visible.




Figure 4.4: Violin plot of DeepEval quality metrics (correctness, relevancy, completeness) by agent configuration. The custom agent exhibits narrower distributions with higher medians across all three metrics.

Figure 4.5 shows the distribution of binary pass/fail outcomes across all queries.



Figure 4.5: Distribution of pass/fail outcomes across all queries for each agent configuration, based on the DeepEval threshold. The custom agent achieves a higher pass rate.

The aggregate score distribution across all DeepEval metrics is shown in Figure 4.6.




Figure 4.6: Aggregate distribution of all DeepEval metric scores. The custom agent’s distribution is more concentrated near the upper end, while the baseline exhibits a heavier low-scoring tail.

The model-level breakdown (Figure 4.7) adds a further dimension.




Figure 4.7: DeepEval quality metrics broken down by underlying LLM. The custom agent outperforms the baseline for every model tested, with the gap narrowing for the most capable model.

The custom agent outperforms the baseline across all three metrics for every model tested, with one exception: for the most capable model in the evaluation, the two configurations converge and the difference becomes negligible. For the smaller and less capable models, the margin is wider. This suggests that structured factsheet context acts as a capability amplifier whose marginal effect is greatest when the base model has limited capacity to compensate for noisy or incomplete retrieval. The implications of this model-dependent effect are discussed further in Chapter 5.

Dataset-level breakdown. Table 4.1 reports per-model scores for both agent configurations across each dataset. Table 4.2 condenses this into mean custom-minus-baseline differences, and Table 4.3 reports the variance of that gap across LLMs.

Table 4.1: Mean completeness, correctness, and relevancy by dataset, model, and agent configuration (Baseline+RAG vs. Custom).

Dataset	Model	Completeness		Correctness		Relevancy	
		Base	Cust	Base	Cust	Base	Cust
THRD-2021-163881	GLM-5	0.880	0.921	0.890	0.928	0.992	1.000
THRD-2021-163881	Qwen3.5-397B	0.937	0.947	0.927	0.944	0.996	1.000
THRD-2021-163881	Gemini 2.5 Flash	0.740	0.853	0.665	0.831	0.996	0.979
THRD-2021-163881	Gemini 2.5 Pro	0.903	0.870	0.894	0.871	1.000	1.000
TOSL-2024-125319	GLM-5	0.992	0.982	0.971	0.956	1.000	1.000
TOSL-2024-125319	Qwen3.5-397B	0.700	0.926	0.727	0.927	0.733	1.000
TOSL-2024-125319	Gemini 2.5 Flash	0.822	0.911	0.718	0.889	0.967	0.980
TOSL-2024-125319	Gemini 2.5 Pro	0.956	0.956	0.911	0.924	0.998	1.000

Table 4.2: Mean difference (custom – baseline) per dataset, in percentage points. Positive values indicate the custom agent scores higher.

Dataset	Completeness	Correctness	Relevancy
THRD-2021-163881	+3.27%	+4.96%	−0.10%
TOSL-2024-125319-MIN	+7.60%	+9.24%	+7.06%

Table 4.3: Variance of the per-model custom–baseline difference per dataset, expressed as a percentage ($\times 100$). Higher variance indicates less uniform gains across LLMs.

Dataset	Completeness	Correctness	Relevancy
THRD-2021-163881	0.38%	0.67%	0.01%
TOSL-2024-125319-MIN	1.20%	1.18%	1.71%

The dataset-level breakdown reveals a consistent asymmetry: the custom agent’s advantage is substantially larger on TOSL-2024-125319 than on THRD-2021-163881 across all three quality dimensions. On THRD-2021-163881, the mean correctness gain is +4.96 percentage points; on TOSL-2024-125319 it more than doubles to +9.24 pp. Completeness follows the same pattern (+3.27 vs. +7.60 pp), while relevancy is effectively neutral on THRD-2021-163881 (−0.10 pp) but shows a clear positive gain on TOSL-2024-125319 (+7.06 pp).

This asymmetry is consistent with the architectural hypothesis. TOSL-2024-125319 is the fully authentic dataset with a larger and more heterogeneous document set; queries on this case require cross-document synthesis that naive RAG is most prone to fail. The factsheet’s structured representation provides the greatest marginal benefit precisely where retrieval-only approaches are most likely to miss facts dispersed across many documents.

The variance table reinforces this interpretation. The variance of the per-model gap is uniformly higher on TOSL-2024-125319 than on THRD-2021-163881 across all metrics. This indicates that the advantage is less uniform across LLMs on the more complex dataset: some models benefit more from the structured context than others, and this interaction is amplified when the case is harder. The notable exception is gemini-2.5-pro on TOSL-2024-125319, where the baseline already achieves high scores, leaving less headroom for improvement.

4.1.2 Query-Type Breakdown

A breakdown of quality metrics by query type provides the most direct empirical evidence for RQ1. As defined in Section 3.3.2, queries are classified as either *factual* (specific, enumerable facts with a bounded correct answer) or *summary* (cross-document synthesis requiring a narrative overview). Summary queries, by their nature, require the agent to synthesise information across multiple documents

and sessions — precisely the capability that naive RAG struggles with most.




Figure 4.8: DeepEval quality metrics broken down by query type (factual vs. summary). The custom agent’s advantage is substantially larger on summary queries, which require cross-document synthesis, than on factual queries.

Table 4.4: Mean correctness, relevancy, and completeness by query type and agent configuration. Completeness is not evaluated on summary queries (see Section 3.3.2).

Query Type	Configuration	Correctness	Relevancy	Completeness
Factual	Baseline+RAG	0.882	0.996	0.910
Factual	Custom	0.901	0.993	0.924
Summary	Baseline+RAG	0.775	0.912	—
Summary	Custom	0.919	0.999	—

The pattern is striking. On factual queries, the custom agent leads by 1.9 percentage points in correctness — a modest but consistent margin. On summary queries, the gap widens to 14.4 percentage points (0.919 vs. 0.775). The same asymmetry holds for relevancy: the baseline’s relevancy on summary queries (0.912) is substantially below its factual query relevancy (0.996), reflecting cases where retrieval fails to surface the distributed context required for a coherent synthesis and the agent produces an incomplete or unfocused response. The custom agent’s relevancy on summary queries (0.999) is effectively at ceiling.

Mann-Whitney U tests confirm that the custom agent’s advantage on summary queries is statistically significant: correctness ($U = 12,813$, $p = 9.3 \times 10^{-3}$) and relevancy ($U = 12,542$, $p = 2.9 \times 10^{-5}$). Since summary queries require cross-document synthesis by definition — they cannot be answered from a single retrieved chunk — this constitutes direct empirical support for RQ1.

4.1.3 BERTScore

BERTScore and Sentence-BERT similarity (defined in Section 3.3.4) serve as a cross-validation signal that is independent of the LLM judge used for the DeepEval metrics. Because BERTScore measures surface-level semantic overlap, it is less sensitive to structured factual completeness than the DeepEval metrics: a response can score well on BERTScore while omitting specific facts that correctness and completeness evaluation would penalise. Agreement between the two metric families therefore provides stronger evidence that observed quality differences reflect genuine architectural effects rather than judge artefacts.




Figure 4.9: Pairplot of BERTScore sub-metrics (precision, recall, F1, and sentence-BERT similarity). Distributions for the two agent configurations largely overlap, indicating comparable surface-level semantic quality.

The distributions have considerable overlap between the two agent types, but the custom agent shows a higher peak with less variation and more consistent surface-level semantic similarity to the reference answers. The baseline has a wider spread with more low-scoring outliers, suggesting it occasionally produces responses less semantically aligned with the reference. The near-identical BERTScore distributions complement rather than contradict the DeepEval findings: both agents produce semantically coherent output, but the custom agent more reliably includes the specific factual content the gold-standard reference requires.




Figure 4.10: BERTScore F1 and sentence-BERT similarity across sessions. Trajectories are flat and closely aligned between configurations, with the custom agent showing marginally narrower confidence intervals.

Across sessions, BERTScore F1 and sentence-BERT similarity are relatively flat and closely aligned for both agent types, with only minor fluctuations. This is expected: surface similarity to reference answers does not change systematically as the case grows, since the phrasing of correct responses does not become more or less reference-like with document volume. The session-level divergence observed in the DeepEval metrics is therefore not an artefact of reference phrasing but reflects genuine differences in factual coverage.

4.1.4 Resource Metrics

Resource metrics (defined in Section 3.3.4) are analysed alongside quality scores to determine whether the custom agent’s architectural overhead is justified by its performance gains.




Figure 4.11: Pairplot of resource metrics (duration in seconds, input tokens, output tokens, total tokens, and LLM call count). Duration and token distributions overlap substantially between configurations; the main visible difference is a more right-skewed LLM call count for the custom agent.

Figure 4.11 shows that the two configurations are broadly similar across duration, input tokens, output tokens, and total tokens — the distributions overlap substantially with no visually clear separation. The most notable difference is in LLM call count: the custom agent’s distribution is more right-skewed, reflecting occasional queries that require additional tool calls to surface structured case facts, while the baseline is more tightly peaked in the lower range (approximately 2–5 calls per query). This is consistent with the custom agent’s richer tool set, which can trigger extra retrieval steps when the factsheet and vector store are both consulted within a single query.




Figure 4.12: Resource metrics across sessions. Token counts and latencies fluctuate without a systematic trend for either configuration, and the two track each other closely throughout.

Figure 4.13 plots inference duration against correctness score for individual queries. The custom agent’s responses cluster more tightly — higher correctness with less spread in duration — while the baseline is more dispersed. The pattern confirms that longer inference time does not produce higher quality: both configurations occupy similar latency ranges, yet the custom agent achieves consistently higher scores within that range.




Figure 4.13: Scatter plot of inference duration (seconds) against correctness score per query. The custom agent’s responses are more tightly concentrated at higher correctness scores. No positive correlation between latency and quality is observed for either configuration.

Unlike the quality metrics, no consistent trend is visible in the resource trajectories over session progression. Token counts and latencies fluctuate without a systematic increase, and the two configurations track each other closely throughout (Figure 4.12). The absence of a growth trend is partly explained by the context compression mechanisms applied to the custom agent: as the factsheet grows, significance filtering and display shortening keep the payload from expanding proportionally.

Cost analysis. Absolute token counts are not directly comparable across datasets with different document volumes and session lengths. The total cost for TOSL-2024-125319 is substantially higher than for THRD-2021-163881, making direct comparison of absolute costs misleading. To enable model-level cost comparisons, each run’s estimated cost (input tokens $\times$ price-per-token + output tokens $\times$ price-per-token) is normalised to a percentile rank within its dataset:

$$\text{rank_cost} = \text{rank}_{\text{pct}}(\text{total_cost}) \quad \text{grouped by dataset}$$

A rank of 1.0 means the most expensive configuration within that dataset; 0.0 means the cheapest. This makes cost comparable across datasets of different scales without requiring absolute price estimates for every provider.

Figures 4.14–4.16 plot cost rank against each quality metric for individual queries.




Figure 4.14: Estimated relative cost rank vs. correctness score per query. The custom agent (orange) occupies the upper-left region for several models, achieving higher correctness at lower or equal relative cost.




Figure 4.15: Estimated relative cost rank vs relevancy score per query. Relevancy is less sensitive to architecture than correctness: both configurations cluster near the top of the relevancy scale regardless of cost, reflecting that on-topic responses are easy to achieve once retrieval succeeds.




Figure 4.16: Estimated relative cost rank vs. completeness score per query. The custom agent dominates the upper portion of the completeness axis at comparable or lower cost, mirroring the correctness pattern.

The correctness and completeness scatter plots reveal a consistent pattern: the custom agent (orange) occupies the upper portion of both axes at comparable or lower relative cost for most models. The most striking case is **gemini-2.5-flash**, where the custom configuration achieves a cost rank of 0.125 (cheapest in the set) while delivering a correctness score of 0.860, versus the baseline’s rank of 0.250 and score of 0.691. Relevancy, by contrast, is less discriminative: both configurations cluster near 1.0 regardless of cost, consistent with the earlier finding that on-topic responses are easy to achieve once retrieval succeeds at all.

Table 4.5 summarises mean cost rank and quality scores per model and agent configuration.

Table 4.5: Mean relative cost rank and quality scores per model and agent configuration. Cost rank is a within-dataset percentile (0 = cheapest, 1 = most expensive).

Model	Agent	Cost rank	Correctness	Relevancy	Completeness
GLM-5	baseline_rag	0.875	0.930	0.996	0.936
GLM-5	custom	1.000	0.942	1.000	0.951
Qwen3.5-397B	baseline_rag	0.438	0.827	0.865	0.818
Qwen3.5-397B	custom	0.438	0.935	1.000	0.936
gemini-2.5-flash	baseline_rag	0.250	0.691	0.981	0.781
gemini-2.5-flash	custom	0.125	0.860	0.980	0.882
gemini-2.5-pro	baseline_rag	0.688	0.902	0.999	0.929
gemini-2.5-pro	custom	0.688	0.898	1.000	0.913

The table highlights the practical trade-off. For three of the four models tested, the custom agent achieves equal or lower cost rank while delivering meaningfully higher correctness and completeness. **gemini-2.5-pro** is the exception: both configurations share the same cost rank and their quality scores are within one percentage point, confirming that the most capable model compensates for the absence of structured context through its own reasoning. The structured factsheet therefore functions most effectively as a cost-efficient capability amplifier for mid-tier models.

Actual vs. Advertised Cost. A further nuance emerges when examining the relationship between advertised input token price and the actual, usage-weighted input cost. The advertised price per token does not fully predict the total cost of a query, because more capable models may use tools more efficiently, requiring fewer tokens overall to arrive at an answer. Figure 4.17 illustrates this by plotting the ranked actual input cost against the advertised price.




Figure 4.17: Actual input cost (ranked) vs. advertised input price per token. The relationship is non-linear, indicating that more capable models can be more token-efficient.

The relationship is notably non-linear. The open-source models, despite having competitive advertised prices, tend to incur a higher relative `actual_input_cost_rank`. Conversely, the closed-source models demonstrate greater token efficiency. `gemini-2.5-pro`, for instance, shows a significant drop in actual cost rank compared to its position on the price axis, performing on par with models that are nominally cheaper per token. `gemini-2.5-flash` also performs well on this metric, delivering a low actual cost. This finding, while secondary to the main research questions, is a valuable insight for practical deployment: the true cost of operating a model is a function of both its price and its underlying capability to reason efficiently and use tools judiciously. The custom agent architecture does not significantly alter this dynamic, as its primary contribution is to the quality and consistency of the context, not the model’s intrinsic reasoning efficiency.

4.2 Statistical Modeling

4.2.1 Mann-Whitney U Test

Hypothesis tests follow the design described in Section 3.3.4. Normality was confirmed to be absent via Shapiro-Wilk ($p < 10^{-18}$ for every metric tested), validating the use of the one-sided Mann-Whitney U test as described in Chapter 3.

Quality metrics. Table 4.6 summarises the results. Correctness and relevancy both show statistically significant advantages for the custom agent at the $\alpha = 0.05$ level. Completeness does not reach significance, though the directional trend is the same.

Table 4.6: Mann-Whitney U test results for quality metrics (one-sided, H_1 : custom > baseline).

Metric	U-statistic	p-value	Result
Correctness	74,780.50	1.03×10^{-2}	Significant ($p < 0.05$)
Relevancy	72,733.00	5.54×10^{-5}	Significant ($p < 0.05$)
Completeness	27,278.50	1.06×10^{-1}	Not significant

Correctness is the most diagnostic quality metric: it is hardest to satisfy, most sensitive to the architectural differences under evaluation, and its significant result ($p = 1.03 \times 10^{-2}$) carries the most interpretive weight.

Relevancy warrants a more cautious reading. It is almost universally high for both agents whenever a response is attempted; the significant difference is primarily driven by cases where the baseline fails to retrieve any relevant documents and produces no substantive answer.

The absence of a significant completeness result ($p = 0.106$) is explained by reduced statistical power rather than an absent effect. As described in Section 3.3.2, completeness is evaluated only on factual queries, which reduces the effective sample to roughly 60% of that available for correctness. The directional trend and mean difference are consistent with correctness — both point in the same direction — but the smaller sample is insufficient to clear the significance threshold. Importantly, Levene’s test finds a significant reduction in completeness *variance* ($p = 0.006$), confirming that the distribution shapes genuinely differ between architectures. The most accurate interpretation is that the evaluation lacks the power to confirm the completeness mean shift at this sample size, not that the shift is absent.

Resource metrics. Neither input token usage ($U = 71,961$, $p = 0.859$) nor inference duration ($U = 69,100$, $p = 0.538$) differs significantly between configurations. Mean input token counts are comparable (custom: $79,319 \pm 82,405$; baseline: $75,395 \pm 78,670$) and mean durations are within one second of each other (custom: $39.8 \text{ s} \pm 27.2 \text{ s}$; baseline: $39.6 \text{ s} \pm 28.9 \text{ s}$). The custom agent achieves its quality improvements without imposing a measurable additional cost in tokens or latency.

4.2.2 Levene’s Test for Equality of Variance

While the Mann-Whitney U test establishes that the custom agent’s scores tend to be stochastically *higher*, it says nothing about the *spread* of those scores. Levene’s test for equality of variance addresses this directly: it tests whether the two configurations exhibit significantly different score variability, without assuming normality. The results are presented in Table 4.7.

Table 4.7: Levene’s test results for equality of variance (custom vs. baseline+RAG). A significant result ($p < 0.05$) indicates the two configurations differ in score variability.

Metric	Levene statistic	p-value	Result
Correctness	18.38	< 0.001	Significant
Relevancy	11.56	0.0007	Significant
Completeness	7.76	0.0056	Significant
BERTScore F1	7.41	0.0066	Significant
Sentence-BERT sim.	16.13	< 0.001	Significant
Duration (s)	0.25	0.620	Not significant
Input tokens	0.02	0.885	Not significant
Self-consistency	0.06	0.810	Not significant

The test returns a significant result for all five quality metrics, including the complementary BERTScore and sentence-BERT similarity where the mean difference is small. This means the custom agent is not only more accurate on average — it is *more consistent*. The baseline’s distribution is significantly wider across all quality dimensions, with a heavier low-scoring tail representing cases of partial or failed retrieval. The custom agent’s reduced variance implies more predictable, reliable responses: fewer failures at the low end and less dependence on whether retrieval happens to surface the right chunks for a given query.

By contrast, variance in inference duration, input token counts, and self-consistency scores does not differ significantly between configurations. The resource overhead of the two agents is therefore comparable not only in expectation but also in spread, and self-consistency variability is similarly matched. The reliability advantage of the custom agent is specific to output quality, not a general property of all metrics.

In a legal setting, this reliability dimension may be as practically significant as the mean improvement. A system that occasionally fails unpredictably is harder to trust and supervise than one that performs steadily, even if its average score is similar.

4.2.3 Self-consistency



Figure 4.18: Distribution of self-consistency scores by agent type. The custom agent shows a higher median and narrower interquartile range.

Self-consistency scores reinforce the pattern seen in the DeepEval metrics visually: the custom agent exhibits a higher median and a narrower interquartile range than the baseline, suggesting more reproducible responses across sessions. However, a one-sided Mann-Whitney U test does not confirm a statistically significant advantage ($U = 7,039$, $p = 0.875$). Similarly, Levene’s test finds no significant difference in self-consistency variance between the two configurations (statistic = 0.06, $p = 0.810$). The descriptive pattern is therefore directionally consistent with the quality findings but does not reach significance, likely due to the smaller number of observations available at the session-group level.

4.3 Summary by Research Question

Table 4.8 summarises the empirical evidence for each research question.

Table 4.8: Research question outcomes based on empirical results.

RQ	Finding	Evidence
Main	Custom agent improves factual accuracy without significant cost increase	Correctness significant ($p=1.03\times 10^{-2}$); token usage and latency statistically equivalent
RQ1	Custom agent significantly outperforms baseline on summary (cross-document) queries; factual query gap is smaller but consistent	Correctness gap +14.4pp on summary queries ($p=9.3\times 10^{-3}$); relevancy gap +8.7pp ($p=2.9\times 10^{-5}$)
RQ2	Indicative evidence from TOSL-2024-125319: gap widens in later sessions; not conclusive with current sample	Session trajectories; wider gap in sessions 3-4 on authentic case only
RQ3	Structured context amplifies smaller model performance; the gap between custom and baseline agents narrows as model capability increases	LLM barplot: gap is widest for smaller models, negligible for the most capable model

Main RQ. The results support an affirmative answer: the custom agent achieves statistically significant improvements in correctness and relevancy over the RAG-only baseline, while consuming a statistically equivalent number of tokens and producing responses within the same latency envelope. The quality gain is not purchased at a proportional resource cost.

RQ1. The query-type breakdown (Section 4.1.2) provides direct empirical support for RQ1. The custom agent’s correctness advantage on summary queries — which require cross-document synthesis by definition — is +14.4 percentage points and statistically significant ($p = 9.3 \times 10^{-3}$). On factual queries, where a single retrieved chunk can often supply the answer, the gap narrows to +1.9pp. The concentration of the advantage in precisely the query type that stresses naive RAG most directly answers RQ1 affirmatively: structured extraction mitigates the cross-document retrieval limitations of chunk-based RAG.

RQ2. Session-level trajectories show that the custom agent maintains quality across all sessions with narrower variance, while the baseline exhibits wider swings and a less stable profile. This supports the conclusion that the structured representation scales with document volume rather than degrading under it.

RQ3. The model-level breakdown provides the clearest evidence for RQ3: the custom agent’s structured context delivers the largest marginal gain for smaller and less capable models, with the gap shrinking to near-zero for the most advanced model tested. This suggests that the factsheet acts as a capability compensator — partially substituting for the reasoning capacity that larger models supply natively.

Chapter 5

Discussion

5.1 Research Question Findings

This section interprets the empirical results in relation to each research question, drawing on the statistical evidence from Chapter 4 and situating the findings within the theoretical context established in Chapter 2.

Main RQ: accuracy improvement without proportional cost. The results support an affirmative answer. The custom agent achieves statistically significant improvements in correctness ($p = 8.5 \times 10^{-3}$) and relevancy ($p = 3.4 \times 10^{-4}$) while consuming a statistically equivalent number of tokens and inference time compared to the RAG-only baseline. The context compression mechanisms — significance filtering and display shortening applied server-side before context injection — are largely responsible for keeping token usage in check as the factsheet grows: rather than injecting the full structured state verbatim, the agent receives a filtered, abbreviated representation tailored to the query context. The combined picture is that schema-driven extraction improves factual accuracy at no measurable additional cost per query.

The non-significant completeness result ($p = 0.106$) should not be read as a negative finding. As described in Section 3.3.2, completeness is only evaluated on factual queries, which reduces the effective sample to roughly 60% of that available for correctness, directly lowering statistical power. The directional trend and mean difference mirror the correctness result; the smaller sample is insufficient to clear the significance threshold. The significant variance reduction in completeness ($p = 0.006$) confirms that the distribution shapes do differ between architectures, and the most accurate interpretation is that the evaluation lacks the power to confirm the mean shift at this sample size. Expanding the evaluation to include more factual queries — particularly those requiring cross-document synthesis — would be expected to bring completeness in line with the correctness result.

Beyond the mean improvement, the variance reduction is arguably the most consistent finding in the evaluation. Levene’s test for equality of variance is significant across all five metrics tested — correctness, relevancy, completeness, BERTScore F1, and sentence-BERT similarity — including the complementary metrics where the mean difference is modest. The baseline agent’s score distributions are significantly wider, with a heavier low-scoring tail that reflects the brittleness of pure retrieval: when the relevant chunks do not rank highly for a given query, performance drops sharply. The custom agent’s structured state provides a reliable fallback that absorbs these retrieval failures, yielding more predictable responses across the board. In a professional legal setting, this consistency advantage may carry more practical weight than the mean improvement alone: a system that occasionally fails unpredictably is harder to trust and integrate into a practitioner’s workflow than one that performs steadily.

RQ1: mitigation of cross-document retrieval limitations. The query-type breakdown provides the clearest evidence in the evaluation. As established in Section 4.1.2, the custom agent’s correctness advantage on summary queries — which require synthesising information across multiple documents — is +14.4 percentage points and statistically significant ($p = 9.3 \times 10^{-3}$). On factual queries, where a single retrieved chunk can often supply the complete answer, the gap narrows to +1.9pp. The

asymmetry is the expected signature of an architecture that specifically addresses cross-document synthesis: it provides the greatest marginal benefit precisely where naive chunk retrieval is most likely to fail.

A theoretical nuance is worth noting: naive RAG is not straightforwardly *volume-sensitive*. Because retrieval always returns a fixed top- K regardless of corpus size, a larger document pool does not mechanically degrade retrieval quality. What RAG genuinely struggles with is *synthesis queries* — questions whose answers are distributed across multiple chunks that cannot all fit within a single top- K result. This is a query-type limitation, not a volume limitation. The query-type analysis directly targets this limitation and finds the expected effect. RQ1 is answered affirmatively: structured extraction mitigates the cross-document synthesis failure of chunk-based RAG, and the effect is statistically significant.

RQ2: stability across growing document volume. Session-level trajectories show that the custom agent maintains consistent quality with narrow variance throughout case progression, while the baseline exhibits wider fluctuations and a less stable profile as sessions accumulate new documents. This is consistent with the design hypothesis: the factsheet absorbs new information incrementally at each session boundary, so the agent’s knowledge of the case becomes richer rather than more fragmented as volume grows. The baseline, by contrast, must retrieve from an ever-larger document pool with no persistent structure to guide retrieval, making its performance increasingly dependent on whether the relevant chunks happen to rank highly in a given semantic search.

RQ3: compensation for reduced model capacity. The model-level breakdown provides the most striking result for RQ3: the performance gap between the custom and baseline agents is largest for the smaller models and shrinks to near-zero for the most capable model tested. For smaller models, the structured factsheet provides the key facts in an organised, pre-digested form, reducing the reasoning burden on the model itself. A more capable model can compensate for noisy or incomplete retrieval through superior inference, making the factsheet less decisive. This model-dependent pattern has a direct practical implication for cost management: a smaller, less expensive model paired with the custom architecture can approximate the performance of a larger, more expensive model on the RAG-only baseline, suggesting a cost-quality substitution that may be valuable in production deployments.

5.2 Value Beyond Performance Metrics

The evaluation results demonstrate a performance advantage and a higher reliability for the custom architecture, but the headline metrics alone understate the practical value of the design. Even in a hypothetical scenario where correctness and completeness scores were identical across configurations, the custom architecture would remain preferable in a production legal setting for reasons that no automatic metric captures.

The factsheet as a product artifact. The baseline agent’s output is ephemeral: it produces a text response and leaves no persistent structured representation of the case behind. The custom agent, by contrast, produces and continuously maintains a structured factsheet as a natural side-effect of normal operation. This factsheet—comprising parties, events, claims, damages, deadlines, and document metadata—is a machine-readable, queryable representation of the case that can be consumed directly by any frontend systems without any additional LLM inference. Concrete applications include timeline visualisations, party graphs, claims dashboards, and deadline trackers, all of which can be rendered from the factsheet with no extra cost.

Human oversight and auditability. The structured state enables a qualitatively different relationship between the lawyer and the system. Rather than trusting a black-box text response, a practitioner can inspect the factsheet directly, verify individual entries against source documents, and correct errors at the element level. If the agent produces an incorrect answer, the source of the error can often be traced to a specific entity entry—a mis-extracted event date or an omitted party—and corrected there, improving all subsequent responses. This kind of targeted, human-in-the-loop correction is not possible with a stateless RAG agent, where errors are buried in the implicit context and must be addressed by re-uploading or re-prompting.

Versioned case understanding. Because the factsheet is updated incrementally at each session boundary, it constitutes a natural audit trail of how the agent’s understanding of the case evolved over time. Earlier versions can be retained and compared, providing a record of when specific facts, parties, or claims entered the agent’s knowledge base. In legal practice, where the timing of knowledge can itself be legally significant, this version history has potential evidentiary and procedural value.

Downstream system integration. The structured state lowers the integration barrier for adjacent legal tools and workflows. A downstream system can query the factsheet via a standard database interface to, for example, check for conflicts of interest against other open cases, trigger calendar entries for extracted deadlines, or populate court filing templates with party and claim data. These integrations require structured, typed data; they cannot be built on top of free-form agent responses without a costly and fragile natural language parsing layer.

Summary. The performance metrics reported in Chapter 4 capture one dimension of the architecture’s value: how accurately the agent answers factual questions across a progression of sessions. They do not capture the value of the structured case representation itself—a persistent, inspectable, and reusable artifact that the custom architecture produces as a consequence of its design. In production deployments with real case volumes and real practitioner workflows, this latent value is likely to be at least as significant as the raw performance advantage.

5.3 Uncertainties and Confounding Factors

The evaluation results support the conclusion that the custom agent outperforms the baseline on correctness and relevancy, but several sources of uncertainty affect the magnitude and generalisability of this finding. These are discussed below.

Question difficulty and scope. The evaluation queries were designed to be factual and document-grounded, but the level of specificity varies considerably across sessions. Some questions have unambiguous single-document answers, while others require synthesising information across multiple documents and timelines. Harder queries amplify the architectural advantage of the custom agent — the factsheet and structured retrieval tools are most useful precisely when the answer requires cross-document reasoning — while easier queries may show little to no difference between architectures. The observed aggregate performance gap therefore depends on the composition of the query set: a test set dominated by simple lookups would understate the custom agent’s advantage, while one focused on complex synthesis would overstate it. The current evaluation mix likely sits somewhere in between.

Dataset size and complexity. Only two datasets are used in the evaluation, and their characteristics differ substantially. The results from TOSL-2024-125319, the larger and more complex real case, show a more pronounced advantage for the custom agent than those from THRD-2021-163881. This is consistent with the design hypothesis: the factsheet-based architecture is most beneficial when the document set is large and the case is complex, because those are the conditions under which naive retrieval is most likely to fail. With only two datasets, however, it is difficult to disentangle the effect of document volume from other case-specific characteristics. A larger benchmark spanning more cases and varying document volumes would provide stronger evidence for the scaling behaviour observed here.

Underlying language model capability. Figure 4.7 shows that the performance gap between the custom and baseline agents narrows as the underlying LLM becomes more capable: for the most advanced model tested (Gemini 2.5 Pro), the two configurations are nearly indistinguishable. This suggests two non-exclusive interpretations. First, a sufficiently capable model may compensate for the structural disadvantage of the baseline by inferring missing context or tolerating retrieval noise more robustly. Second, the evaluation queries and datasets may not be sufficiently demanding to expose the limits of a powerful model: the questions can largely be answered from the retrieved chunks even without the accumulated factsheet context. Conversely, the performance gap is most pronounced for the smaller models, for which structured, pre-digested context provides the greatest marginal benefit.

This implies that the value of the custom architecture is not fixed but model-dependent: it acts as a capability amplifier that matters most when the underlying model has limited reasoning capacity or context-handling ability.

Stochastic variability. LLM outputs are inherently stochastic. Each evaluation set was run three times to reduce the impact of this variability, and the non-parametric Mann-Whitney U test is robust to the non-normal, left-skewed score distributions observed. Nevertheless, the total number of unique query-session combinations is constrained by the cost and time of running large-scale multi-model evaluations, and the reported p-values should be interpreted with awareness of the limited sample size. The significance of correctness ($p = 8.5 \times 10^{-3}$) and relevancy ($p = 3.4 \times 10^{-4}$) results are robust, but the borderline result for completeness ($p = 8.3 \times 10^{-2}$) should be treated cautiously.

5.4 Overall Assessment

Taken together, the findings of this thesis present a coherent picture. The core claim — that structured extraction into a persistent case representation improves factual accuracy without proportional cost — is supported by statistically significant evidence on the most diagnostic metrics, consistent directional trends across all others, and a robust reduction in output variance that holds across every metric tested. The uncertainties identified above are real but do not undermine this conclusion: they define the boundaries within which it applies and point toward the conditions under which the advantage grows or shrinks.

What the aggregate statistics alone do not convey is that the value of the architecture extends well beyond the evaluation setting. The factsheet is not merely a retrieval aid — it is a persistent, inspectable, and continuously updated representation of the case that practitioners can verify, correct, and build upon. In a professional legal context, where trust in a tool is built incrementally through transparent and auditable behaviour, this structural property may prove more consequential than any marginal gain in benchmark scores.

The picture that emerges is one of complementarity rather than substitution. Structured extraction does not replace retrieval — it fills the gap that retrieval cannot address on its own: the coherent, cross-document case overview that no similarity score can assemble from isolated chunks. The two mechanisms together — a persistent structured state and a semantic vector store — provide coverage that neither offers alone. This complementarity is the architectural insight this thesis contributes, and it is one that holds regardless of the specific models, datasets, or evaluation metrics used to measure it.

Chapter 6

Conclusion

This thesis investigated whether schema-driven cross-document extraction into a structured case representation — a *FactSheet* — improves the factual accuracy of an agentic RAG system in multi-document legal settings, and whether it does so without incurring proportional increases in query-time cost. The motivation was grounded in a fundamental limitation of chunk-based retrieval: when the answer to a legal query is distributed across multiple documents, no amount of semantic similarity tuning can make an isolated-chunk retriever assemble a coherent synthesis. The proposed approach addresses this by extracting typed entities — parties, events, claims, damages, deadlines — from all case documents at ingestion time and maintaining them as a persistent, relational representation that supplements retrieval for case-level queries.

6.1 Answers to Research Questions

Main RQ. Schema-driven extraction into a structured case representation does improve factual accuracy: the custom agent achieves statistically significant improvements in correctness ($p = 8.5 \times 10^{-3}$) and relevancy ($p = 3.4 \times 10^{-4}$) relative to the RAG-only baseline. Token consumption and inference latency are statistically equivalent between the two architectures, confirming that the quality gain is not purchased at a proportional resource cost. The main research question is answered affirmatively.

RQ1. The query-type breakdown provides direct empirical support for RQ1. Summary queries — which require synthesising information across multiple documents — show a +14.4 percentage-point correctness advantage for the custom agent ($p = 9.3 \times 10^{-3}$), whereas factual queries, where a single retrieved chunk can often supply the complete answer, show only a +1.9pp gap. This asymmetry is the expected signature of an architecture that specifically addresses cross-document synthesis: it provides the greatest marginal benefit precisely where naive chunk retrieval is most likely to fail. A theoretical nuance applies: naive RAG is not straightforwardly volume-sensitive, since retrieval always returns a fixed top- K regardless of corpus size. The genuine limitation is query-type specific — synthesis queries whose answers are distributed across multiple chunks — and the query-type analysis directly targets this limitation. RQ1 is answered affirmatively.

RQ2. Session-level trajectories show that the custom agent maintains stable quality with narrow variance across all sessions as document volume grows. The baseline, by contrast, exhibits wider fluctuations throughout case progression. The structured representation scales gracefully with document volume because new information is absorbed incrementally into the factsheet at each session boundary rather than added to an ever-larger unstructured retrieval pool. This pattern is consistent with the design hypothesis, but with only two datasets it is difficult to isolate the effect of document volume from other case-specific characteristics. RQ2 is indicative but not conclusively confirmed; a larger benchmark spanning more cases and varying document volumes would be required for a firm answer.

RQ3. The model-level breakdown reveals that the performance gap between the custom and baseline agents is largest for smaller, less capable models and narrows to near-zero for the most advanced model tested. Structured, pre-digested context reduces the reasoning burden placed on the underlying model,

enabling smaller models to approximate the performance of larger ones on the RAG-only baseline. This constitutes a practically significant cost-quality substitution: a cheaper model paired with the custom architecture can match a more expensive model operating on unstructured retrieval alone. RQ3 is answered affirmatively, with the caveat that the effect is model-dependent and diminishes as base model capability increases.

6.2 Additional Finding: Variance Reduction

Beyond the hypothesised mean improvement, the evaluation reveals a consistent and statistically robust reduction in score variance for the custom agent. Levene’s test for equality of variance is significant across all five metrics tested — correctness, relevancy, completeness, BERTScore F1, and sentence-BERT similarity — including the complementary metrics where mean differences are modest. This finding was not pre-specified as a research question; it emerged from the distributional analysis of the results.

The interpretation is straightforward: pure retrieval is brittle. When relevant chunks happen not to rank highly for a given query, the baseline agent’s performance drops sharply, producing the heavy low-scoring tail visible across all its metric distributions. The custom agent’s structured state provides a stable knowledge base that does not depend on retrieval succeeding for any individual query, dampening these failures and yielding more predictable responses.

In a legal deployment context, this reliability dimension may carry more practical weight than the mean improvement. A practitioner integrating an AI assistant into their workflow needs to be able to predict when the system will and will not perform well. A system with lower variance is easier to trust, easier to supervise, and easier to integrate into structured legal processes than one with comparable average accuracy but unpredictable failure modes.

6.3 Contributions

This thesis makes three contributions:

1. **System design.** A schema-driven initialization and incremental update pipeline for structured legal case state management, implemented as a LangGraph `StateGraph` with fan-out parallel processing, PostgreSQL relational storage, and configurable context compression.
2. **Evaluation framework.** A multi-model evaluation setup comparing a structured-state agent against a RAG-only baseline across two legal datasets, multiple session stages, and three repeated runs per configuration, using a combination of semantic quality metrics (DeepEval correctness, relevancy, completeness), surface similarity metrics (BERTScore), consistency metrics, and resource consumption measures.
3. **Empirical evidence.** Quantitative evidence that structured extraction improves factual accuracy in multi-document legal settings without proportional cost increase, with the effect being strongest on complex cases, stable across session progression, and most pronounced for smaller models.

6.4 Limitations and Future Work

The evaluation in this thesis has several limitations that point toward avenues for future research.

Dataset and Generalisability. The evaluation rests on two datasets, which limits the generalisability of the findings. A larger benchmark spanning more cases and document volumes would provide stronger evidence for the scaling behaviour observed here, and would better distinguish the effect of document volume from other case-specific characteristics. A potential weakness is that the datasets were developed in conjunction with this thesis, which introduces a risk that their design was implicitly influenced by the research questions, potentially limiting their independence as a neutral benchmark.

A more fundamental limitation, however, is the scarcity of publicly available, high-quality datasets suitable for benchmarking legal agents in realistic, long-duration scenarios. The datasets used, while

adequate for the research questions posed, represent a simplified proxy for the complexity of a genuine legal case that contains hundreds of sessions with thousands of individual queries. The evaluation was necessarily limited to a sequence of sessions within a static, pre-existing document corpus. A crucial direction for future work is therefore the development of high-quality, publicly available datasets designed specifically for full-scale, real-world scenario testing. Such benchmarks are essential for the field to move beyond simple, short-form question-answering formats and begin to test agent capabilities in a manner that mirrors actual legal practice. These datasets should be designed to test agents across time, presenting new information and evolving case dynamics in a structured way that challenges an agent’s ability to maintain a coherent, long-term understanding of a case. Creating such a benchmark is a substantial research project in its own right, arguably the subject of a separate thesis. It would require careful curation of case materials, the design of a longitudinal evaluation protocol, and the definition of tasks that probe for deep synthesis and long-range reasoning rather than simple fact retrieval. Making such datasets publicly available would additionally enable other researchers to benchmark alternative architectures under comparable conditions, accelerating progress across the field rather than limiting reproducibility to groups with access to proprietary case files.

Scope of Legal Reasoning. The evaluation queries were deliberately restricted to factual, document-grounded questions to enable objective grading. This scope is appropriate for establishing a controlled baseline but does not cover the full range of legal reasoning tasks — in particular, normative and predictive questions where the benefit of structured extraction may differ. Extending the evaluation to cover multi-step legal reasoning, conflicting evidence reconciliation, and temporal argument construction is a natural next step.

Schema Flexibility. The structured representation in this thesis uses a fixed entity schema (parties, events, claims, damages, deadlines). For cases with different legal structures, such as regulatory compliance matters or transactional disputes, the schema may require adaptation. An interesting direction is schema induction: learning the relevant entity types from the case documents themselves rather than relying on a predefined legal ontology.

Extraction Quality Evaluation. The evaluation in this thesis measures downstream output quality — whether the agent’s answers are correct, relevant, and complete — but does not assess the intermediate quality of the extracted factsheet itself. A factsheet element that is incorrectly extracted, hallucinated, or incompletely attributed goes undetected unless it demonstrably degrades a downstream answer. This is a meaningful blind spot: unlike a failed retrieval, which simply omits a relevant chunk, an incorrect factsheet entry is injected directly and authoritatively into the model’s context, where it may actively mislead the agent across many subsequent queries.

A dedicated extraction evaluation would require expert-annotated ground truth factsheets for each case, against which extracted elements could be assessed for precision, recall, and source attribution accuracy. Evaluation should cover the full range of entity types — parties, events, claims, damages, deadlines — and account for partial matches, scope errors (where the correct entity type is identified but the extracted value is incomplete), and hallucinations with no grounding in the source documents. It would also be valuable to study how extraction quality varies by document type: contracts, email correspondence, and financial reports present different linguistic registers and information densities, and the extraction pipeline may perform unevenly across them. Establishing this link between extraction quality and downstream answer quality would clarify how much of the observed performance gap is attributable to the structured representation itself versus the precision with which it is populated.

Cost-Quality Optimisation and Model Combination. The RQ3 result points toward a practically relevant optimisation problem: given a target quality level and a budget constraint, what is the optimal combination of model size and extraction investment? The evidence suggests this trade-off exists and is exploitable, but characterising it precisely requires a more systematic sweep over model and architecture configurations than was feasible within the scope of this thesis.

A natural extension is to decouple the model used for extraction from the model used for inference. In the current system, the same model handles both pipeline extraction and query-time reasoning. An alternative configuration would use a more capable model for the extraction pipeline — where high-quality structured output matters most and cost is amortised over many queries — while using a

smaller, cheaper model at query time where volume is high. This combination is practically attractive and directly testable with data already collected in this thesis: since the factsheets have been extracted and stored, inference-only runs on those projects can be executed without re-running the extraction pipeline, making the marginal cost of the additional comparison low.

Reliability as a First-Class Research Question. The variance reduction observed across all five metrics is a post-hoc finding and was not designed into the evaluation. A natural follow-up study would treat reliability as a first-class research question: formally hypothesising that structured extraction reduces output variance, designing the evaluation around variance-sensitive query sets (particularly retrieval-failure-prone queries), and testing across a broader range of cases and document volumes. Given the practical significance of predictability in legal settings, this direction warrants dedicated investigation.

Advanced Retrieval over Structured Elements. The current implementation queries the fact-sheet through simple SQL lookups filtered by date range and significance level. The richness of the extracted representation — typed entities with source references, timestamps, significance scores, and relational links — supports substantially more sophisticated retrieval strategies. Graph traversal across entity relationships, cross-referencing events with the documents that substantiate them, and ranking elements by their relevance to a specific query rather than by static significance are all unexplored. The factsheet as implemented is a data foundation whose retrieval layer has not been fully exploited. It is found in other related works that the quality in the retrieval strategy yields larger gains than sophisticated extraction and write strategies. [17] Of course this presumes that the underlying data is of sufficient quality, which the structured extraction supports by providing a clean, typed, and source-linked representation. But the retrieval layer itself is a critical component that can be further optimised to maximise the value of the structured data.

Agentic Reasoning and Task Queuing. The current agent executes a single reasoning pass per query. Adding a task queue that allows the agent to decompose complex queries into ordered sub-tasks — each potentially using a specialised node, toolset, and system prompt — would support multi-step legal reasoning that the current flat architecture cannot perform. This is analogous to chain-of-thought planning at the agent level: rather than producing an answer in one shot, the agent could plan, retrieve, verify, and synthesise across multiple coordinated steps.

Active Context Monitoring. The current implementation does not actively monitor context utilisation at runtime. Systems such as MemGPT [16] demonstrate that tool-mediated relief — notifying the model when its context budget is near exhaustion and providing tools to archive or compact in-context content — can substantially extend coherent long-horizon behaviour. Implementing equivalent monitoring within the LangGraph agent, with automatic compaction of conversation history and selective eviction of lower-priority factsheet elements, would be a direct extension of the current architecture.

Background and Asynchronous Tooling. The pipeline already runs offline, but the agent itself is purely reactive: it only acts when a query arrives. A natural extension is a set of background tools that run during idle periods — monitoring connected document repositories and inboxes for new material, triggering incremental factsheet updates on arrival, or proactively flagging newly detected deadlines and conflicts to the practitioner. This would move the system from reactive assistant to proactive case monitor.

Graph Database and GraphRAG. The relational PostgreSQL schema captures entity types and source links but does not natively represent or traverse relationships between entities. Migrating the structured representation to a graph database would enable direct traversal of entity relationships — for example, following a chain from a claim to the events that substantiate it to the documents that evidence those events. Combined with graph-native retrieval, this would address the cross-document reasoning limitation more completely than the current hybrid approach and bring the system closer to a full domain-specific GraphRAG architecture.

Bibliography

- [1] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need, 2023.
- [2] Cheng-Ping Hsieh, Simeng Sun, Samuel Kriman, Shantanu Acharya, Dima Rekeshe, Fei Jia, Yang Zhang, and Boris Ginsburg. Ruler: What’s the real context size of your long-context language models? *arXiv preprint arXiv:2404.06654*, 2024.
- [3] Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranajape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts, 2023.
- [4] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks, 2021.
- [5] Yunfan Gao, Yun Xiong, Yijie Zhong, Yuxi Bi, Ming Xue, and Haofen Wang. Synergizing rag and reasoning: A systematic review, 2025.
- [6] Yixuan Tang and Yi Yang. Multihop-RAG: Benchmarking retrieval-augmented generation for multi-hop queries. In *First Conference on Language Modeling*, 2024.
- [7] Jintao Liang, Gang Su, Huifeng Lin, You Wu, Rui Zhao, and Ziyue Li. Reasoning rag via system 1 or system 2: A survey on reasoning agentic retrieval-augmented generation for industry challenges, 2025.
- [8] Pengcheng Jiang, Siru Ouyang, Yizhu Jiao, Ming Zhong, Runchu Tian, and Jiawei Han. Retrieval and structuring augmented generation with large language models. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*, KDD ’25, page 6032–6042. ACM, August 2025.
- [9] LangChain. Langgraph, 2024. Accessed: 2025-11-26.
- [10] Microsoft. Semantic kernel python: Context management, 2024. Accessed: 2025-11-26.
- [11] Ilias Chalkidis, Manos Fergadiotis, Prodromos Malakasiotis, Nikolaos Aletras, and Ion Androutsopoulos. LEGAL-BERT: The muppets straight out of law school. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 2898–2904. Association for Computational Linguistics, 2020.
- [12] Neel Guha, Julian Nyarko, Daniel E. Ho, Christopher Ré, Adam Chilton, Aditya Narayana, Alex Chohlas-Wood, Austin Peters, Brandon Waldon, Daniel N. Rockmore, Diego Zambrano, Dmitry Talisman, Enam Hoque, Faiz Surani, Frank Fagan, Galit Sarfaty, Gregory M. Dickinson, Haggai Porat, Jason Hegland, Jessica Wu, Joe Nudell, Joel Niklaus, John Nay, Jonathan H. Choi, Kevin Tobia, Margaret Hagan, Megan Ma, Michael Livermore, Nikon Rasumov-Rahe, Nils Holzenberger, Noam Kolt, Peter Henderson, Sean Rehaag, Sharad Goel, Shang Gao, Spencer Williams, Sunny Gandhi, Tom Zur, Varun Iyer, and Zehua Li. LegalBench: A collaboratively built benchmark for measuring legal reasoning in large language models, 2023.
- [13] Antoine Louis, Gijs van Dijck, and Gerasimos Spanakis. Interpretable long-form legal question answering with retrieval-augmented large language models, 2023.

- [14] Hao Liu, Zhengren Wang, Xi Chen, Zhiyu Li, Feiyu Xiong, Qinhan Yu, and Wentao Zhang. Hoprag: Multi-hop reasoning for logic-aware retrieval-augmented generation. *arXiv preprint arXiv:2502.12442*, 2025.
- [15] Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, Dasha Metropolitansky, Robert Osazuwa Ness, and Jonathan Larson. From local to global: A graph rag approach to query-focused summarization, 2025.
- [16] Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. Memgpt: Towards llms as operating systems, 2024.
- [17] Boqin Yuan, Yue Su, and Kun Yao. Diagnosing retrieval vs. utilization bottlenecks in llm agent memory, 2026.
- [18] Parth Sarthi, Salman Abdullah, Aditi Tuli, Shubh Khanna, Anna Goldie, and Christopher D. Manning. Raptor: Recursive abstractive processing for tree-organized retrieval, 2024.
- [19] Sebastián Ramírez. Fastapi, 2018. Accessed: 2025-11-26.
- [20] Supabase. Supabase, 2020.
- [21] Zejiang Shen, Kyle Lo, Lauren Yu, Nathan Dahlberg, Margo Schlanger, and Doug Downey. Multi-lexsum: Real-world summaries of civil rights lawsuits at multiple granularities, 2022.
- [22] Haitao Li, Jiaying Ye, Yiran Hu, Jia Chen, Qingyao Ai, Yueyue Wu, Junjie Chen, Yifan Chen, Cheng Luo, Quan Zhou, and Yiqun Liu. Casegen: A benchmark for multi-stage legal case documents generation, 2025.
- [23] Michael J. Ryan, Danmei Xu, Chris Nivera, and Daniel Campos. Enronqa: Towards personalized rag over private documents, 2025.
- [24] Neel Guha, Julian Nyarko, Daniel E. Ho, Christopher Ré, Adam Chilton, Aditya Narayana, Alex Chohlas-Wood, Austin Peters, Brandon Waldon, Daniel N. Rockmore, Diego Zambrano, Dmitry Talisman, Enam Hoque, Faiz Surani, Frank Fagan, Galit Sarfaty, Gregory M. Dickinson, Haggai Porat, Jason Hegland, Jessica Wu, Joe Nudell, Joel Niklaus, John Nay, Jonathan H. Choi, Kevin Tobia, Margaret Hagan, Megan Ma, Michael Livermore, Nikon Rasumov-Rahe, Nils Holzenberger, Noam Kolt, Peter Henderson, Sean Rehaag, Sharad Goel, Shang Gao, Spencer Williams, Sunny Gandhi, Tom Zur, Varun Iyer, and Zehua Li. Legalbench: A collaboratively built benchmark for measuring legal reasoning in large language models, 2023.
- [25] Ilias Chalkidis, Abhik Jana, Dirk Hartung, Michael Bommarito, Ion Androutsopoulos, Daniel Martin Katz, and Nikolaos Aletras. Lexglue: A benchmark dataset for legal language understanding in english, 2022.
- [26] Confident AI. Deepeval: The llm evaluation framework, 2023.
- [27] Qwen Team. Qwen3.5: Towards native multimodal agents, February 2026.
- [28] Together AI. Together ai. [urlhttps://together.ai/](https://together.ai/), 2024. Accessed: 2026-04-08.
- [29] Gemini Team, Rohan Anil, Sebastian Borgeaud, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M. Dai, Anja Hauth, Katie Millican, David Silver, Melvin Johnson, Ioannis Antonoglou, Julian Schrittwieser, Amelia Glaese, Jilin Chen, Emily Pitler, Timothy Lili-crap, Angeliki Lazaridou, Orhan Firat, James Molloy, Michael Isard, Paul R. Barham, Tom Hen-nigan, Benjamin Lee, Fabio Viola, Malcolm Reynolds, Yuanzhong Xu, Ryan Doherty, Eli Collins, Clemens Meyer, Eliza Rutherford, Erica Moreira, Kareem Ayoub, Megha Goel, Jack Krawczyk, Cosmo Du, Ed Chi, Heng-Tze Cheng, Eric Ni, Purvi Shah, Patrick Kane, Betty Chan, Man-aal Faruqui, Aliaksei Severyn, Hanzhao Lin, YaGuang Li, Yong Cheng, Abe Ittycheriah, Mahdis Mahdieh, Mia Chen, Pei Sun, Dustin Tran, Sumit Bagri, Balaji Lakshminarayanan, Jeremiah Liu, Andras Orban, Fabian Gura, Hao Zhou, Xinying Song, Aurelien Boffy, Harish Ganapathy,

Steven Zheng, HyunJeong Choe, Ágoston Weisz, Tao Zhu, Yifeng Lu, Siddharth Gopal, Jarrod Kahn, Maciej Kula, Jeff Pitman, Rushin Shah, Emanuel Taropa, Majd Al Merey, Martin Baeuml, Zhifeng Chen, Laurent El Shafey, Yujing Zhang, Olcan Sercinoglu, George Tucker, Enrique Piqueras, Maxim Krikun, Iain Barr, Nikolay Savinov, Ivo Danihelka, Becca Roelofs, Anaïs White, Anders Andreassen, Tamara von Glehn, Lakshman Yagati, Mehran Kazemi, Lucas Gonzalez, Misha Khalman, Jakub Sygnowski, Alexandre Frechette, Charlotte Smith, Laura Culp, Lev Proleev, Yi Luan, Xi Chen, James Lottes, Nathan Schucher, Federico Lebron, Alban Rrustemi, Natalie Clay, Phil Crone, Tomas Kocisky, Jeffrey Zhao, Bartek Perz, Dian Yu, Heidi Howard, Adam Bloniarz, Jack W. Rae, Han Lu, Laurent Sifre, Marcello Maggioni, Fred Alcober, Dan Garrette, Megan Barnes, Shantanu Thakoor, Jacob Austin, Gabriel Barth-Maron, William Wong, Rishabh Joshi, Rahma Chaabouni, Deeni Fatiha, Arun Ahuja, Gaurav Singh Tomar, Evan Senter, Martin Chadwick, Ilya Kornakov, Nithya Attaluri, Iñaki Iturrate, Ruibo Liu, Yunxuan Li, Sarah Cogan, Jeremy Chen, Chao Jia, Chenjie Gu, Qiao Zhang, Jordan Grimstad, Ale Jakse Hartman, Xavier Garcia, Thanumalayan Sankaranarayana Pillai, Jacob Devlin, Michael Laskin, Diego de Las Casas, Dasha Valter, Connie Tao, Lorenzo Blanco, Adrià Puigdomènech Badia, David Reitter, Mianna Chen, Jenny Brennan, Clara Rivera, Sergey Brin, Shariq Iqbal, Gabriela Surita, Jane Labanowski, Abhi Rao, Stephanie Winkler, Emilio Parisotto, Yiming Gu, Kate Olszewska, Ravi Addanki, Antoine Miech, Annie Louis, Denis Teplyashin, Geoff Brown, Elliot Catt, Jan Balaguer, Jackie Xiang, Pidong Wang, Zoe Ashwood, Anton Briukhov, Albert Webson, Sanjay Ganapathy, Smit Sanghavi, Ajay Kannan, Ming-Wei Chang, Axel Stjerngren, Josip Djolonga, Yuting Sun, Ankur Bapna, Matthew Aitchison, Pedram Pejman, Henryk Michalewski, Tianhe Yu, Cindy Wang, Juliette Love, Junwhan Ahn, Dawn Bloxwich, Kehang Han, Peter Humphreys, Thibault Sellam, James Bradbury, Varun Godbole, Sina Samangooei, Bogdan Damoc, Alex Kaskasoli, Sébastien M. R. Arnold, Vijay Vasudevan, Shubham Agrawal, Jason Riesa, Dmitry Lepikhin, Richard Tanburn, Srivatsan Srinivasan, Hyeontaek Lim, Sarah Hodgkinson, Pranav Shyam, Johan Ferret, Steven Hand, Ankush Garg, Tom Le Paine, Jian Li, Yujia Li, Minh Giang, Alexander Neitz, Zaheer Abbas, Sarah York, Machel Reid, Elizabeth Cole, Aakanksha Chowdhery, Dipanjan Das, Dominika Rogozińska, Vitaliy Nikolaev, Pablo Sprechmann, Zachary Nado, Lukas Zilka, Flavien Prost, Luheng He, Marianne Monteiro, Gaurav Mishra, Chris Welty, Josh Newlan, Dawei Jia, Miltiadis Allamanis, Clara Huiyi Hu, Raoul de Liedekerke, Justin Gilmer, Carl Saroufim, Shruti Rijhwani, Shaobo Hou, Disha Shrivastava, Anirudh Baddepudi, Alex Goldin, Adnan Ozturel, Albin Cassirer, Yunhan Xu, Daniel Sohn, Devendra Sachan, Reinald Kim Amplayo, Craig Swanson, Dessie Petrova, Shashi Narayan, Arthur Guez, Siddhartha Brahma, Jessica Landon, Miteyan Patel, Ruizhe Zhao, Kevin Vilella, Luyu Wang, Wenhao Jia, Matthew Rahtz, Mai Giménez, Legg Yeung, James Keeling, Petko Georgiev, Diana Mincu, Boxi Wu, Salem Haykal, Rachel Saputro, Kiran Vodrahalli, James Qin, Zeynep Cankara, Abhanshu Sharma, Nick Fernando, Will Hawkins, Behnam Neyshabur, Solomon Kim, Adrian Hutter, Priyanka Agrawal, Alex Castro-Ros, George van den Driessche, Tao Wang, Fan Yang, Shuo yiin Chang, Paul Komarek, Ross McIlroy, Mario Lučić, Guodong Zhang, Wael Farhan, Michael Sharman, Paul Natsev, Paul Michel, Yamini Bansal, Siyuan Qiao, Kris Cao, Siamak Shakeri, Christina Butterfield, Justin Chung, Paul Kishan Rubenstein, Shivani Agrawal, Arthur Mensch, Kedar Soparkar, Karel Lenc, Timothy Chung, Aedan Pope, Loren Maggiore, Jackie Kay, Priya Jhakra, Shibo Wang, Joshua Maynez, Mary Phuong, Taylor Tobin, Andrea Tacchetti, Maja Trebacz, Kevin Robinson, Yash Katariya, Sebastian Riedel, Paige Bailey, Kefan Xiao, Nimesh Ghelani, Lora Aroyo, Ambrose Slone, Neil Houlsby, Xuehan Xiong, Zhen Yang, Elena Gribovskaya, Jonas Adler, Mateo Wirth, Lisa Lee, Music Li, Thais Kagohara, Jay Pavagadhi, Sophie Bridgers, Anna Bortsova, Sanjay Ghemawat, Zafarali Ahmed, Tianqi Liu, Richard Powell, Vijay Bolina, Mariko Iinuma, Polina Zablotskaia, James Besley, Da-Woon Chung, Timothy Dozat, Ramona Comanescu, Xiance Si, Jeremy Greer, Guolong Su, Martin Polacek, Raphaël Lopez Kaufman, Simon Tokumine, Hexiang Hu, Elena Buchatskaya, Yingjie Miao, Mohamed Elhawaty, Aditya Siddhant, Nenad Tomasev, Jinwei Xing, Christina Greer, Helen Miller, Shereen Ashraf, Aurko Roy, Zizhao Zhang, Ada Ma, Angelos Filos, Milos Besta, Rory Blevins, Ted Klimenko, Chih-Kuan Yeh, Soravit Changpinyo, Jiaqi Mu, Oscar Chang, Mantas Pajarskas, Carrie Muir, Vered Cohen, Charline Le Lan, Krishna Haridasan, Amit Marathe, Steven Hansen, Sholto Douglas, Rajkumar Samuel, Mingqiu Wang, Sophia Austin, Chang Lan, Jiepu Jiang, Justin Chiu, Jaime Alonso Lorenzo, Lars Lowe Sjösund, Sébastien Cevey, Zach Gleicher, Thi Avrahami, Anudhyan Boral, Hansa Srinivasan,

Vittorio Selo, Rhys May, Konstantinos Aisopos, Léonard Hussenot, Livio Baldini Soares, Kate Baumli, Michael B. Chang, Adrià Recasens, Ben Caine, Alexander Pritzel, Filip Pavetic, Fabio Pardo, Anita Gergely, Justin Frye, Vinay Ramasesh, Dan Horgan, Kartikeya Badola, Nora Kassner, Subhrajit Roy, Ethan Dyer, Víctor Campos Campos, Alex Tomala, Yunhao Tang, Dalia El Badawy, Elspeth White, Basil Mustafa, Oran Lang, Abhishek Jindal, Sharad Vikram, Zhitao Gong, Sergi Caelles, Ross Hemsley, Gregory Thornton, Fangxiaoyu Feng, Wojciech Stokowiec, Ce Zheng, Phoebe Thacker, Çağlar Ünlü, Zhishuai Zhang, Mohammad Saleh, James Svensson, Max Bileschi, Piyush Patil, Ankesh Anand, Roman Ring, Katerina Tsihla, Arpi Vezer, Marco Selvi, Toby Shevlane, Mikel Rodriguez, Tom Kwiatkowski, Samira Daruki, Keran Rong, Allan Dafoe, Nicholas FitzGerald, Keren Gu-Lemberg, Mina Khan, Lisa Anne Hendricks, Marie Pellet, Vladimir Feinberg, James Cobon-Kerr, Tara Sainath, Maribeth Rauh, Sayed Hadi Hashemi, Richard Ives, Yana Hasson, Eric Noland, Yuan Cao, Nathan Byrd, Le Hou, Qingze Wang, Thibault Sottiaux, Michela Paganini, Jean-Baptiste Lespiau, Alexandre Moufarek, Samer Hassan, Kaushik Shivakumar, Joost van Amersfoort, Amol Mandhane, Pratik Joshi, Anirudh Goyal, Matthew Tung, Andrew Brock, Hannah Sheahan, Vedant Misra, Cheng Li, Nemanja Rakićević, Mostafa Dehghani, Fangyu Liu, Sid Mittal, Junhyuk Oh, Seb Noury, Eren Sezener, Fantine Huot, Matthew Lamm, Nicola De Cao, Charlie Chen, Sidharth Mudgal, Romina Stella, Kevin Brooks, Gautam Vasudevan, Chenxi Liu, Mainak Chaitin, Nivedita Melinkeri, Aaron Cohen, Venus Wang, Kristie Seymore, Sergey Zubkov, Rahul Goel, Summer Yue, Sai Krishnakumaran, Brian Albert, Nate Hurley, Motoki Sano, Anhad Mohananey, Jonah Joughin, Egor Filonov, Tomasz Kępa, Yomna Eldawy, Jiawern Lim, Rahul Rishi, Shirin Badiezhadegan, Taylor Bos, Jerry Chang, Sanil Jain, Sri Gayatri Sundara Padmanabhan, Subha Puttagunta, Kalpesh Krishna, Leslie Baker, Norbert Kalb, Vamsi Bedapudi, Adam Kurzrok, Shuntong Lei, Anthony Yu, Oren Litvin, Xiang Zhou, Zhichun Wu, Sam Sobell, Andrea Siciliano, Alan Papir, Robby Neale, Jonas Bragagnolo, Tej Toor, Tina Chen, Valentin Anklin, Feiran Wang, Richie Feng, Milad Gholami, Kevin Ling, Lijuan Liu, Jules Walter, Hamid Moghaddam, Arun Kishore, Jakub Adamek, Tyler Mercado, Jonathan Mallinson, Siddhinita Wandekar, Stephen Cagle, Eran Ofek, Guillermo Garrido, Clemens Lombriser, Maksim Mukha, Botu Sun, Hafeezul Rahman Mohammad, Josip Matak, Yadi Qian, Vikas Peswani, Pawel Janus, Quan Yuan, Leif Schelin, Oana David, Ankur Garg, Yifan He, Oleksii Duzhyi, Anton Älgmyr, Timothée Lottaz, Qi Li, Vikas Yadav, Luyao Xu, Alex Chinien, Rakesh Shivanna, Aleksandr Chuklin, Josie Li, Carrie Spadine, Travis Wolfe, Kareem Mohamed, Subhabrata Das, Zihang Dai, Kyle He, Daniel von Dincklage, Shyam Upadhyay, Akanksha Maurya, Luyan Chi, Sebastian Krause, Khalid Salama, Pam G Rabinovitch, Pavan Kumar Reddy M, Aarush Selvan, Mikhail Dektiarev, Golnaz Ghiasi, Erdem Guven, Himanshu Gupta, Boyi Liu, Deepak Sharma, Idan Heimlich Shtacher, Shachi Paul, Oscar Akerlund, François-Xavier Aubet, Terry Huang, Chen Zhu, Eric Zhu, Elico Teixeira, Matthew Fritze, Francesco Bertolini, Liana-Eleonora Marinescu, Martin Böhle, Dominik Paulus, Khyatti Gupta, Tejas Latkar, Max Chang, Jason Sanders, Roopa Wilson, Xuewei Wu, Yi-Xuan Tan, Lam Nguyen Thiet, Tulsee Doshi, Sid Lall, Swaroop Mishra, Wanming Chen, Thang Luong, Seth Benjamin, Jasmine Lee, Ewa Andrejczuk, Dominik Rabiej, Vipul Ranjan, Krzysztof Styrz, Pengcheng Yin, Jon Simon, Malcolm Rose Harriott, Mudit Bansal, Alexei Robsky, Geoff Bacon, David Greene, Daniil Mirylenka, Chen Zhou, Obaid Sarvana, Abhimanyu Goyal, Samuel Andermatt, Patrick Siegler, Ben Horn, Assaf Israel, Francesco Pongetti, Chih-Wei "Louis" Chen, Marco Selvatici, Pedro Silva, Kathie Wang, Jackson Tolins, Kelvin Guu, Roey Yogev, Xiaochen Cai, Alessandro Agostini, Maulik Shah, Hung Nguyen, Noah Ó Donnaile, Sébastien Pereira, Linda Friso, Adam Stambler, Adam Kurzrok, Chenkai Kuang, Yan Romanikhin, Mark Geller, ZJ Yan, Kane Jang, Cheng-Chun Lee, Wojciech Fica, Eric Malmi, Qijun Tan, Dan Banica, Daniel Balle, Ryan Pham, Yanping Huang, Diana Avram, Hongzhi Shi, Jasjit Singh, Chris Hidey, Niharika Ahuja, Pranab Saxena, Dan Dooley, Srividya Pranavi Potharaju, Eileen O'Neill, Anand Gokulchandran, Ryan Foley, Kai Zhao, Mike Dusenberry, Yuan Liu, Pulkit Mehta, Ragha Kotikalapudi, Chalence Safranek-Shrader, Andrew Goodman, Joshua Kessinger, Eran Globen, Prateek Kolhar, Chris Gorgolewski, Ali Ibrahim, Yang Song, Ali Eichenbaum, Thomas Brovelli, Sahitya Potluri, Preethi Lahoti, Cip Baetu, Ali Ghorbani, Charles Chen, Andy Crawford, Shalini Pal, Mukund Sridhar, Petru Gurita, Asier Mujika, Igor Petrovski, Pierre-Louis Cedoz, Chenmei Li, Shiyuan Chen, Niccolò Dal Santo, Siddharth Goyal, Jitesh Punjabi, Karthik Kappaganthu, Chester Kwak, Pallavi LV, Sarmishta Velury, Himadri Choudhury, Jamie Hall, Premal Shah, Ricardo Figueira, Matt Thomas, Minjie Lu, Ting

Zhou, Chintu Kumar, Thomas Jurdi, Sharat Chikkerur, Yenai Ma, Adams Yu, Soo Kwak, Victor Ähdel, Sujeevan Rajayogam, Travis Choma, Fei Liu, Aditya Barua, Colin Ji, Ji Ho Park, Vincent Hellendoorn, Alex Bailey, Taylan Bilal, Huanjie Zhou, Mehrdad Khatir, Charles Sutton, Wojciech Rzadkowski, Fiona Macintosh, Roopali Vij, Konstantin Shagin, Paul Medina, Chen Liang, Jin-jing Zhou, Pararth Shah, Yingying Bi, Attila Dankovics, Shipra Banga, Sabine Lehmann, Marissa Bredesen, Zifan Lin, John Eric Hoffmann, Jonathan Lai, Raynald Chung, Kai Yang, Nihal Balani, Arthur Bražinskas, Andrei Sozanschi, Matthew Hayes, Héctor Fernández Alcalde, Peter Makarov, Will Chen, Antonio Stella, Liselotte Snijders, Michael Mandl, Ante Kärman, Paweł Nowak, Xinyi Wu, Alex Dyck, Krishnan Vaidyanathan, Raghavender R, Jessica Mallet, Mitch Rudominer, Eric Johnston, Sushil Mittal, Akhil Udathu, Janara Christensen, Vishal Verma, Zach Irving, Andreas Santucci, Gamaleldin Elsayed, Elnaz Davoodi, Marin Georgiev, Ian Tenney, Nan Hua, Geoffrey Cideron, Edouard Leurent, Mahmoud Alnahlawi, Ionut Georgescu, Nan Wei, Ivy Zheng, Dylan Scandinaro, Heinrich Jiang, Jasper Snoek, Mukund Sundararajan, Xuezhi Wang, Zack Ontiveros, Itay Karo, Jeremy Cole, Vinu Rajashekhar, Lara Tumeh, Eyal Ben-David, Rishub Jain, Jonathan Uesato, Romina Datta, Oskar Bunyan, Shimu Wu, John Zhang, Piotr Stanczyk, Ye Zhang, David Steiner, Subhajit Naskar, Michael Azzam, Matthew Johnson, Adam Paszke, Chung-Cheng Chiu, Jaume Sanchez Elias, Afroz Mohiuddin, Faizan Muhammad, Jin Miao, Andrew Lee, Nino Vieillard, Jane Park, Jiageng Zhang, Jeff Stanway, Drew Garmon, Abhijit Karmarkar, Zhe Dong, Jong Lee, Aviral Kumar, Luowei Zhou, Jonathan Evens, William Isaac, Geoffrey Irving, Edward Loper, Michael Fink, Isha Arkatkar, Nanxin Chen, Izhak Shafran, Ivan Petrychenko, Zhe Chen, Johnson Jia, Anselm Levskaya, Zhenkai Zhu, Peter Grabowski, Yu Mao, Alberto Magni, Kaisheng Yao, Javier Snaider, Norman Casagrande, Evan Palmer, Paul Suganthan, Alfonso Castaño, Irene Giannoumis, Wooyeol Kim, Mikołaj Rybiński, Ashwin Sreevatsa, Jennifer Prendki, David Soergel, Adrian Goedeckemeyer, Willi Gierke, Mohsen Jafari, Meenu Gaba, Jeremy Wiesner, Diana Gage Wright, Yawen Wei, Harsha Vashisht, Yana Kulizhskaya, Jay Hoover, Maigo Le, Lu Li, Chimezie Iwuanyanwu, Lu Liu, Kevin Ramirez, Andrey Khorlin, Albert Cui, Tian LIN, Marcus Wu, Ricardo Aguilar, Keith Pallo, Abhishek Chakladar, Ginger Perng, Elena Allica Abellan, Mingyang Zhang, Ishita Dasgupta, Nate Kushman, Ivo Penchev, Alena Repina, Xihui Wu, Tom van der Weide, Priya Ponnappalli, Caroline Kaplan, Jiri Simsa, Shuangfeng Li, Olivier Dousse, Fan Yang, Jeff Piper, Nathan Ie, Rama Pasumarthi, Nathan Lintz, Anitha Vijayakumar, Daniel Andor, Pedro Valenzuela, Minnie Lui, Cosmin Paduraru, Daiyi Peng, Katherine Lee, Shuyuan Zhang, Somer Greene, Duc Dung Nguyen, Paula Kurylowicz, Cassidy Hardin, Lucas Dixon, Lili Janzer, Kiam Choo, Ziqiang Feng, Biao Zhang, Achintya Singhal, Dayou Du, Dan McKinnon, Natasha Antropova, Tolga Bolukbasi, Orgad Keller, David Reid, Daniel Finchelstein, Maria Abi Raad, Remi Crocker, Peter Hawkins, Robert Dadashi, Colin Gaffney, Ken Franko, Anna Bulanova, Rémi Leblond, Shirley Chung, Harry Askham, Luis C. Cobo, Kelvin Xu, Felix Fischer, Jun Xu, Christina Sorokin, Chris Alberti, Chu-Cheng Lin, Colin Evans, Alek Dimitriev, Hannah Forbes, Dylan Banarse, Zora Tung, Mark Omernick, Colton Bishop, Rachel Sterneck, Rohan Jain, Jiawei Xia, Ehsan Amid, Francesco Piccinno, Xingyu Wang, Praseem Banzal, Daniel J. Mankowitz, Alex Polozov, Victoria Krakovna, Sasha Brown, MohammadHossein Bateni, Dennis Duan, Vlad Firoiu, Meghana Thotakuri, Tom Natan, Matthieu Geist, Ser tan Girgin, Hui Li, Jiayu Ye, Ofir Roval, Reiko Tojo, Michael Kwong, James Lee-Thorp, Christopher Yew, Danila Sinopalnikov, Sabela Ramos, John Mellor, Abhishek Sharma, Kathy Wu, David Miller, Nicolas Sonnerat, Denis Vnukov, Rory Greig, Jennifer Beattie, Emily Caveness, Libin Bai, Julian Eisenschlos, Alex Korchemniy, Tomy Tsai, Mimi Jasarevic, Weize Kong, Phuong Dao, Zeyu Zheng, Frederick Liu, Fan Yang, Rui Zhu, Tian Huey Teh, Jason Sanmiya, Evgeny Gladchenko, Nejc Trdin, Daniel Toyama, Evan Rosen, Sasan Tavakkol, Linting Xue, Chen Elkind, Oliver Woodman, John Carpenter, George Papamakarios, Rupert Kemp, Sushant Kafle, Tanya Grunina, Rishika Sinha, Alice Talbert, Diane Wu, Denese Owusu-Afriyie, Cosmo Du, Chloe Thornton, Jordi Pont-Tuset, Pradyumna Narayana, Jing Li, Saaber Fatehi, John Wieting, Omar Ajmeri, Benigno Uria, Yeongil Ko, Laura Knight, Amélie Héliou, Ning Niu, Shane Gu, Chenxi Pang, Yeqing Li, Nir Levine, Ariel Stolovich, Rebeca Santamaria-Fernandez, Sonam Goenka, Wenny Yustalim, Robin Strudel, Ali Elqursh, Charlie Deck, Hyo Lee, Zonglin Li, Kyle Levin, Raphael Hoffmann, Dan Holtmann-Rice, Olivier Bachem, Sho Arora, Christy Koh, Soheil Hassas Yeganeh, Siim Põder, Mukarram Tariq, Yanhua Sun, Lucian Ionita, Mojtaba Seyedhosseini, Pouya Tafti, Zhiyu Liu, Anmol Gulati, Jasmine Liu, Xinyu Ye, Bart Chrzasczcz, Lily Wang, Nikhil Sethi, Tianrun Li, Ben Brown, Shreya Singh, Wei Fan, Aaron

Parisi, Joe Stanton, Vinod Koverkathu, Christopher A. Choquette-Choo, Yunjie Li, TJ Lu, Abe Ittycheriah, Prakash Shroff, Mani Varadarajan, Sanaz Bahargam, Rob Willoughby, David Gaddy, Guillaume Desjardins, Marco Cornero, Brona Robenek, Bhavishya Mittal, Ben Albrecht, Ashish Shenoy, Fedor Moiseev, Henrik Jacobsson, Alireza Ghaffarkhah, Morgane Rivière, Alanna Walton, Clément Crepy, Alicia Parrish, Zongwei Zhou, Clement Farabet, Carey Radebaugh, Praveen Srinivasan, Claudia van der Salm, Andreas Fildjeland, Salvatore Scellato, Eri Latorre-Chimoto, Hanna Klimczak-Plucińska, David Bridson, Dario de Cesare, Tom Hudson, Piermaria Mendolicchio, Lexi Walker, Alex Morris, Matthew Mauger, Alexey Guseynov, Alison Reid, Seth Odoom, Lucia Locher, Victor Cotruta, Madhavi Yenugula, Dominik Grewe, Anastasia Petrushkina, Tom Duerig, Antonio Sanchez, Steve Yadlowsky, Amy Shen, Amir Globerson, Lynette Webb, Sahil Dua, Dong Li, Surya Bhupatiraju, Dan Hurt, Haroon Qureshi, Ananth Agarwal, Tomer Shani, Matan Eyal, Anuj Khare, Shreyas Rammohan Belle, Lei Wang, Chetan Tekur, Mihir Sanjay Kale, Jinliang Wei, Ruoxin Sang, Brennan Saeta, Tyler Liechty, Yi Sun, Yao Zhao, Stephan Lee, Pandu Nayak, Doug Fritz, Manish Reddy Vuyyuru, John Aslanides, Nidhi Vyas, Martin Wicke, Xiao Ma, Evgenii Eltyshhev, Nina Martin, Hardie Cate, James Manyika, Keyvan Amiri, Yelin Kim, Xi Xiong, Kai Kang, Florian Luisier, Nilesh Tripuraneni, David Madras, Mandy Guo, Austin Waters, Oliver Wang, Joshua Ainslie, Jason Baldridge, Han Zhang, Garima Pruthi, Jakob Bauer, Feng Yang, Riham Mansour, Jason Gelman, Yang Xu, George Polovets, Ji Liu, Honglong Cai, Warren Chen, Xiang-Hai Sheng, Emily Xue, Sherjil Ozair, Christof Angermueller, Xiaowei Li, Anoop Sinha, Weiren Wang, Julia Wiesinger, Emmanouil Koukoumidis, Yuan Tian, Anand Iyer, Madhu Gurumurthy, Mark Goldenson, Parashar Shah, MK Blake, Hongkun Yu, Anthony Urbanowicz, Jennimaria Palomaki, Chrisantha Fernando, Ken Durden, Harsh Mehta, Nikola Momchev, Elahe Rahimtoroghi, Maria Georgaki, Amit Raul, Sebastian Ruder, Morgan Redshaw, Jinhyuk Lee, Denny Zhou, Komal Jalan, Dinghua Li, Blake Hechtman, Parker Schuh, Milad Nasr, Kieran Milan, Vladimir Mikulik, Juliana Franco, Tim Green, Nam Nguyen, Joe Kelley, Aroma Mahendru, Andrea Hu, Joshua Howland, Ben Vargas, Jeffrey Hui, Kshitij Bansal, Vikram Rao, Rakesh Ghiya, Emma Wang, Ke Ye, Jean Michel Sarr, Melanie Moranski Preston, Madeleine Elish, Steve Li, Aakash Kaku, Jigar Gupta, Ice Pasupat, Da-Cheng Juan, Milan Someswar, Tejvi M., Xinyun Chen, Aida Amini, Alex Fabrikant, Eric Chu, Xuanyi Dong, Amruta Muthal, Senaka Butthipitiya, Sarthak Jauhari, Nan Hua, Urvashi Khandelwal, Ayal Hitron, Jie Ren, Larissa Rinaldi, Shahar Drath, Avigail Dabush, Nan-Jiang Jiang, Harshal Godhia, Uli Sachs, Anthony Chen, Yicheng Fan, Hagai Taitelbaum, Hila Noga, Zhuyun Dai, James Wang, Chen Liang, Jenny Hamer, Chun-Sung Ferng, Chenel Elkind, Aviel Atias, Paulina Lee, Vít Listík, Mathias Carlen, Jan van de Kerkhof, Marcin Pikus, Krunoslav Zaher, Paul Müller, Sasha Zykova, Richard Stefanec, Vitaly Gatsko, Christoph Hirschschall, Ashwin Sethi, Xingyu Federico Xu, Chetan Ahuja, Beth Tsai, Anca Stefanoiu, Bo Feng, Keshav Dhandhanania, Manish Katyal, Akshay Gupta, Atharva Parulekar, Divya Pitta, Jing Zhao, Vivaan Bhatia, Yashodha Bhavnani, Omar Alhadlaq, Xiaolin Li, Peter Danenberg, Dennis Tu, Alex Pine, Vera Filippova, Abhipso Ghosh, Ben Limonchik, Bhargava Urala, Chaitanya Krishna Lanka, Derik Clive, Yi Sun, Edward Li, Hao Wu, Kevin Hongtongsak, Ianna Li, Kalind Thakkar, Kuanysh Omarov, Kushal Majmudar, Michael Alverson, Michael Kucharski, Mohak Patel, Mudit Jain, Maksim Zabelin, Paolo Pelagatti, Rohan Kohli, Saurabh Kumar, Joseph Kim, Swetha Sankar, Vineet Shah, Lakshmi Ramachandruni, Xiangkai Zeng, Ben Bariach, Laura Weidinger, Tu Vu, Alek Andreev, Antoine He, Kevin Hui, Sheleem Kashem, Amar Subramanya, Sissie Hsiao, Demis Hassabis, Koray Kavukcuoglu, Adam Sadovsky, Quoc Le, Trevor Strohman, Yonghui Wu, Slav Petrov, Jeffrey Dean, and Oriol Vinyals. Gemini: A family of highly capable multimodal models, 2025.

- [30] Qwen Team. Qwen3 technical report, 2025.
- [31] GLM-5-Team, :, Aohan Zeng, Xin Lv, Zhenyu Hou, Zhengxiao Du, Qinkai Zheng, Bin Chen, Da Yin, Chendi Ge, Chenghua Huang, Chengxing Xie, Chenzheng Zhu, Congfeng Yin, Cunxiang Wang, Gengzheng Pan, Hao Zeng, Haoke Zhang, Haoran Wang, Huilong Chen, Jiajie Zhang, Jian Jiao, Jiaqi Guo, Jingsen Wang, Jingzhao Du, Jinzhu Wu, Kedong Wang, Lei Li, Lin Fan, Lucen Zhong, Mingdao Liu, Mingming Zhao, Pengfan Du, Qian Dong, Rui Lu, Shuang-Li, Shulin Cao, Song Liu, Ting Jiang, Xiaodong Chen, Xiaohan Zhang, Xuancheng Huang, Xuezhen Dong, Yabo Xu, Yao Wei, Yifan An, Yilin Niu, Yitong Zhu, Yuanhao Wen, Yukuo Cen, Yushi Bai, Zhongpei Qiao, Zihan Wang, Zikang Wang, Zilin Zhu, Ziqiang Liu, Zixuan Li, Bojie Wang, Bosi

- Wen, Can Huang, Changpeng Cai, Chao Yu, Chen Li, Chengwei Hu, Chenhui Zhang, Dan Zhang, Daoyan Lin, Dayong Yang, Di Wang, Ding Ai, Erle Zhu, Fangzhou Yi, Feiyu Chen, Guohong Wen, Hailong Sun, Haisha Zhao, Haiyi Hu, Hanchen Zhang, Hanrui Liu, Hanyu Zhang, Hao Peng, Hao Tai, Haobo Zhang, He Liu, Hongwei Wang, Hongxi Yan, Hongyu Ge, Huan Liu, Huanpeng Chu, Jia'ni Zhao, Jiachen Wang, Jiajing Zhao, Jiamin Ren, Jiapeng Wang, Jiaxin Zhang, Jiayi Gui, Jiayue Zhao, Jijie Li, Jing An, Jing Li, Jingwei Yuan, Jinhua Du, Jinxin Liu, Junkai Zhi, Junwen Duan, Kaiyue Zhou, Kangjian Wei, Ke Wang, Keyun Luo, Laiqiang Zhang, Leigang Sha, Liang Xu, Lindong Wu, Lintao Ding, Lu Chen, Minghao Li, Nianyi Lin, Pan Ta, Qiang Zou, Rongjun Song, Ruiqi Yang, Shangqing Tu, Shangtong Yang, Shaoxiang Wu, Shengyan Zhang, Shijie Li, Shuang Li, Shuyi Fan, Wei Qin, Wei Tian, Weining Zhang, Wenbo Yu, Wenjie Liang, Xiang Kuang, Xiangmeng Cheng, Xiangyang Li, Xiaoquan Yan, Xiaowei Hu, Xiaoying Ling, Xing Fan, Xingye Xia, Xinyuan Zhang, Xinze Zhang, Xirui Pan, Xu Zou, Xunkai Zhang, Yadi Liu, Yandong Wu, Yanfu Li, Yidong Wang, Yifan Zhu, Yijun Tan, Yilin Zhou, Yiming Pan, Ying Zhang, Yinpei Su, Yipeng Geng, Yong Yan, Yonglin Tan, Yuean Bi, Yuhan Shen, Yuhao Yang, Yujiang Li, Yunan Liu, Yunqing Wang, Yuntao Li, Yurong Wu, Yutao Zhang, Yuxi Duan, Yuxuan Zhang, Zezhen Liu, Zhengtao Jiang, Zhenhe Yan, Zheyu Zhang, Zhixiang Wei, Zhuo Chen, Zhuoer Feng, Zijun Yao, Ziwei Chai, Ziyuan Wang, Zuzhou Zhang, Bin Xu, Minlie Huang, Hongning Wang, Juanzi Li, Yuxiao Dong, and Jie Tang. Glm-5: from vibe coding to agentic engineering, 2026.
- [32] Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. Bertscore: Evaluating text generation with bert, 2020.
- [33] Howard Levene. Robust tests for equality of variances. In Ingram Olkin, Harold Hotelling, et al., editors, *Contributions to Probability and Statistics: Essays in Honor of Harold Hotelling*, pages 278–292. Stanford University Press, Stanford, CA, 1960.
- [34] Norwegian University of Life Sciences (NMBU). Guidelines for the use of artificial intelligence (AI) at NMBU, 2024. Accessed: 2026-04-07.